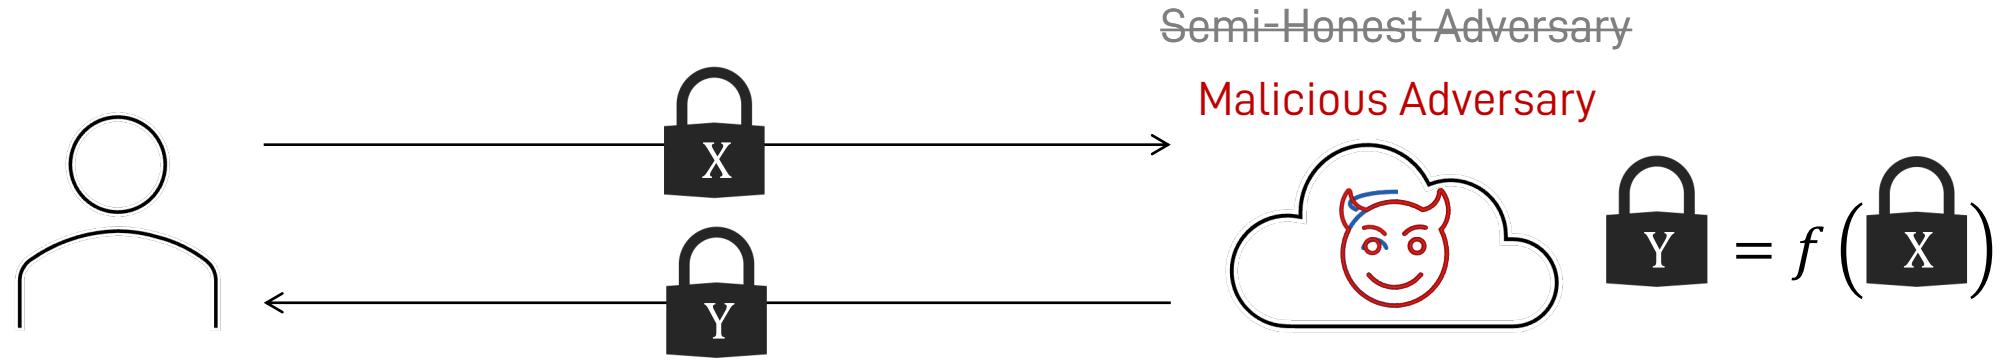


# Verifiable Fully Homomorphic Encryption

Christian Knabenhans\*, Alexander Viand\*, Anwar Hithnawi



# A stronger adversarial model for FHE



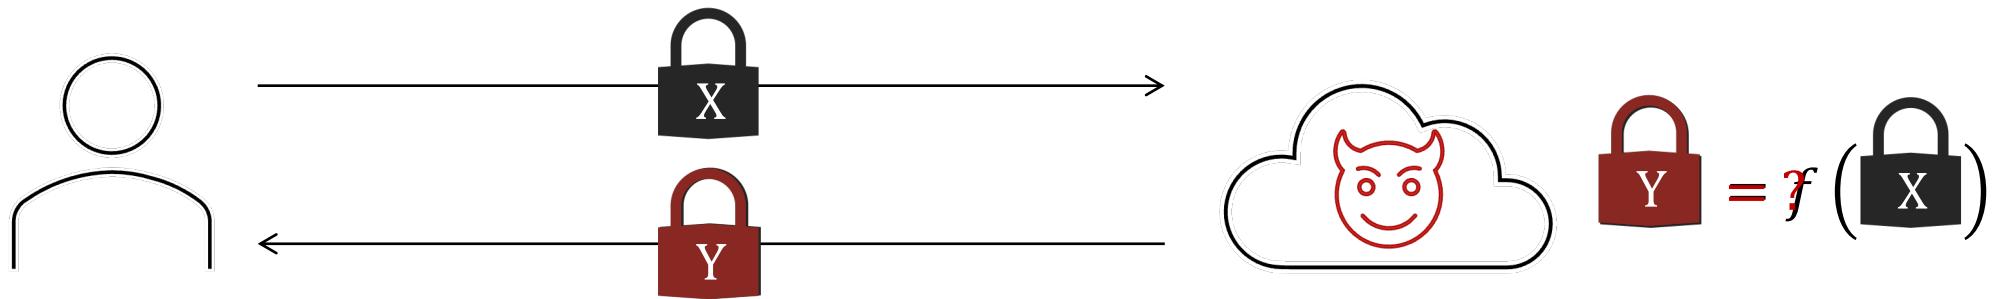
No trusted server

Breaches

Programming errors

Monetary incentive

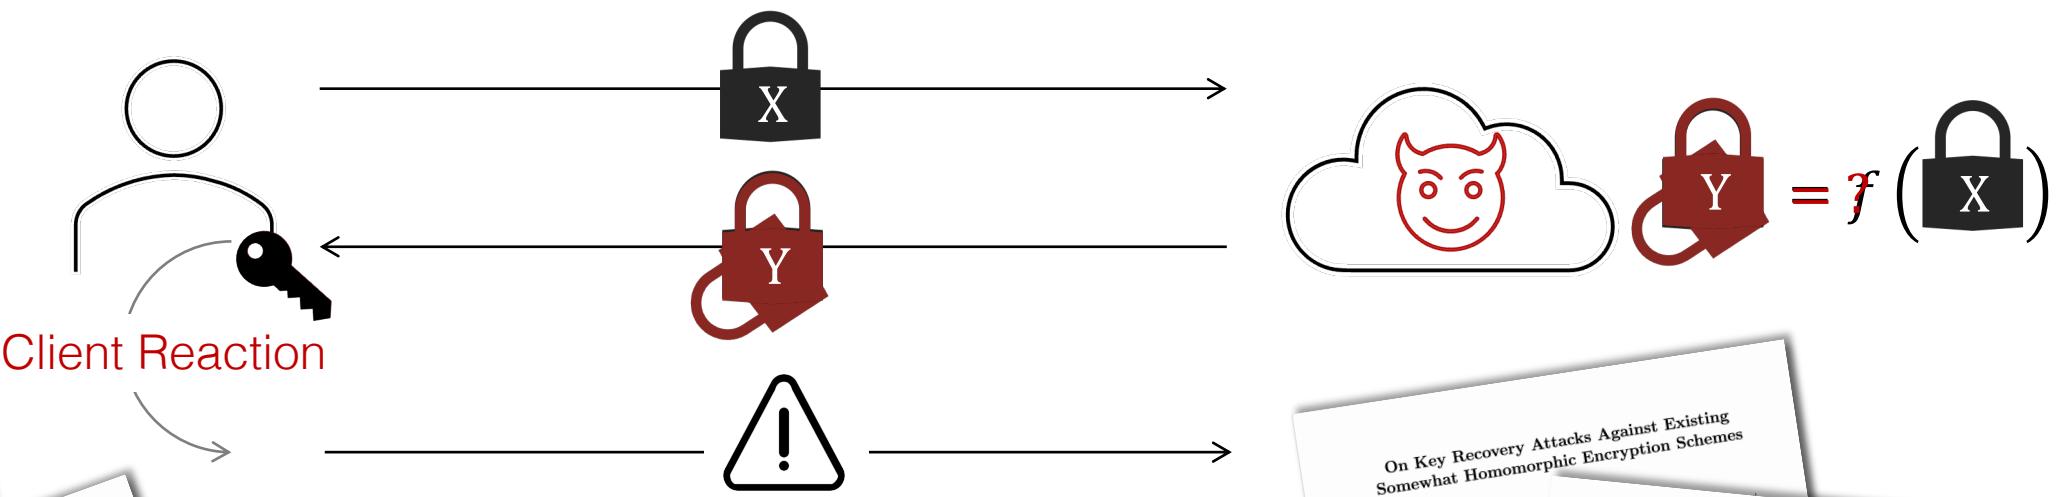
# FHE correctness in malicious setting



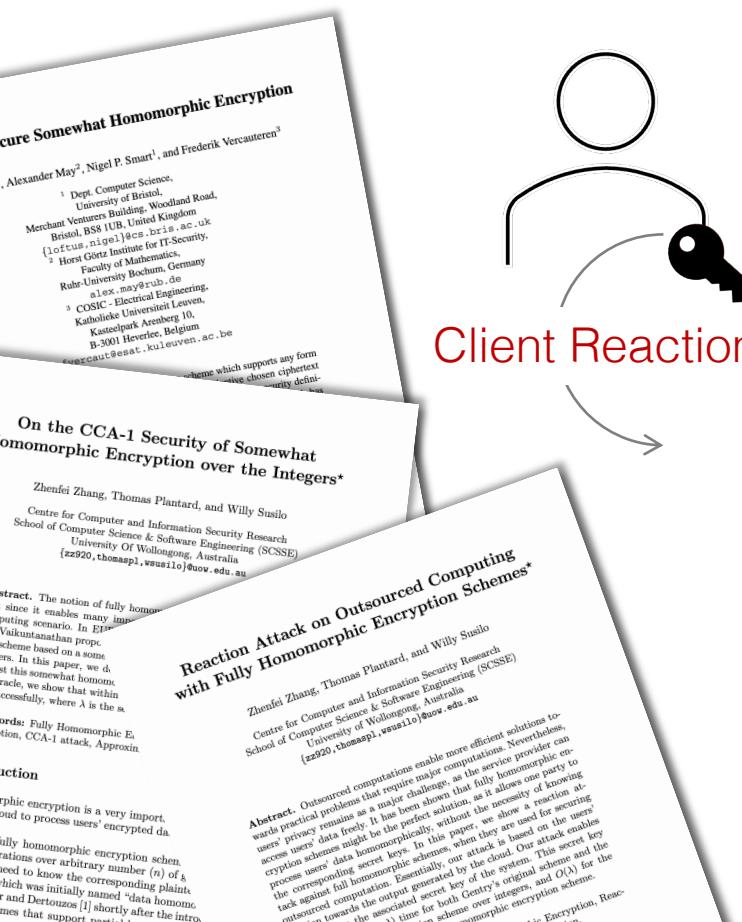
No robust correctness

# FHE confidentiality in malicious setting

Dec( X )  
Dec( X' )  
...



No robust confidentiality



Ilaria Chillotti<sup>1</sup>, Nicolas Gama<sup>2,3</sup>, Louis Goubin<sup>1</sup>  
<sup>1</sup> Laboratoire de Mathématiques de Versailles, UVSQ, CNRS, Université Paris-Saclay, 78035 Versailles, France  
<sup>2</sup> Inpher, Lausanne, Switzerland

---

# Verifiable Fully Homomorphic Encryption

Alexander Viand\*, Christian Knabenhans\*, Anwar Hithnawi

ETH Zurich

**Abstract**—Fully Homomorphic Encryption (FHE) is seeing increasing real-world deployment to protect data in use by allowing computation over encrypted data. However, the same malleability that enables homomorphic computations also raises integrity issues, which have so far been mostly overlooked. While FHE’s lack of integrity has obvious implications for correctness, it also has severe implications for confidentiality: a malicious server can leverage the lack of integrity to carry out interactive key-recovery attacks. As a result, virtually all FHE schemes and applications assume an honest-but-curious server who does not deviate from the protocol. In practice, however, this assumption is insufficient for a wide range of deployment scenarios. While there has been work that aims to address this gap, these have remained isolated efforts considering only aspects of the overall problem and fail to fully address the needs and characteristics of modern FHE schemes and applications. In this paper, we analyze existing FHE integrity approaches, present attacks that exploit gaps in prior work, and propose a new notion for maliciously-secure verifiable FHE. We then instantiate this new notion with a range of techniques, analyzing them and evaluating their performance in a range of different settings. We highlight their potential but also show where future work on tailored integrity solutions for FHE is still required.

## 1. Introduction

Fully Homomorphic Encryption (FHE), which enables computation on encrypted data, has the potential to

**Honest-but-Curious Assumption.** Historically, the FHE research community has extensively made use of the assumption that the server running an FHE application would be honest-but-curious, rather than actively malicious [13]–[16]. This assumption may be reasonable in some deployment scenarios (e.g., when FHE is used only to ensure regulatory compliance or when dealing with trusted institutions cooperating on their own data). However, the necessity to trust the server to this extent is very limiting to the scope of application scenarios, since a violation of the assumption threatens not only correctness but also confidentiality. In addition, even otherwise trusted parties can be compromised by malicious third parties, exposing this attack surface. While FHE protects against passive attacks, a malicious or compromised server taking part in an FHE application can leverage this to undermine data confidentiality. A class of attacks known as *key-recovery attacks* exploits the interactive nature of real-world deployments to construct (partial) decryption oracles. These exploit the fact that a server can craft a ciphertext that fails to decrypt correctly for certain secret keys, using the client’s reaction or lack thereof as an oracle. Practical key-recovery attacks have been developed for all major FHE schemes [8]–[12]. Therefore, there is an urgent need to strengthen FHE to maintain strong guarantees in the context of these attacks.

**Existing FHE Integrity Approaches.** In order to remediate these attacks, a line of research has emerged that constructs more robust FHE schemes [17]–[24] that achieve indistinguishability against chosen ciphertext attacks (IND-CCA1).

# What do we mean with "FHE"?

Formal FHE

Early FHE

Modern FHE



An FHE scheme is a tuple of PPT algorithms ( $KGen, Enc, Eval, Dec$ ) such that ...

Fully Homomorphic Encryption from Ring-LWE and Security for Key Dependent Messages

Zvika Brakerski<sup>1</sup> and Vinod Vaikuntanathan<sup>2</sup>

BV scheme

B/FV, BGV, TFHE, FHEW, CKKS



homenc/HElib  
HElib is an open-source software library that implements homomorphic encryption. It supports the BGV scheme with bootstrapping and the Approximate...

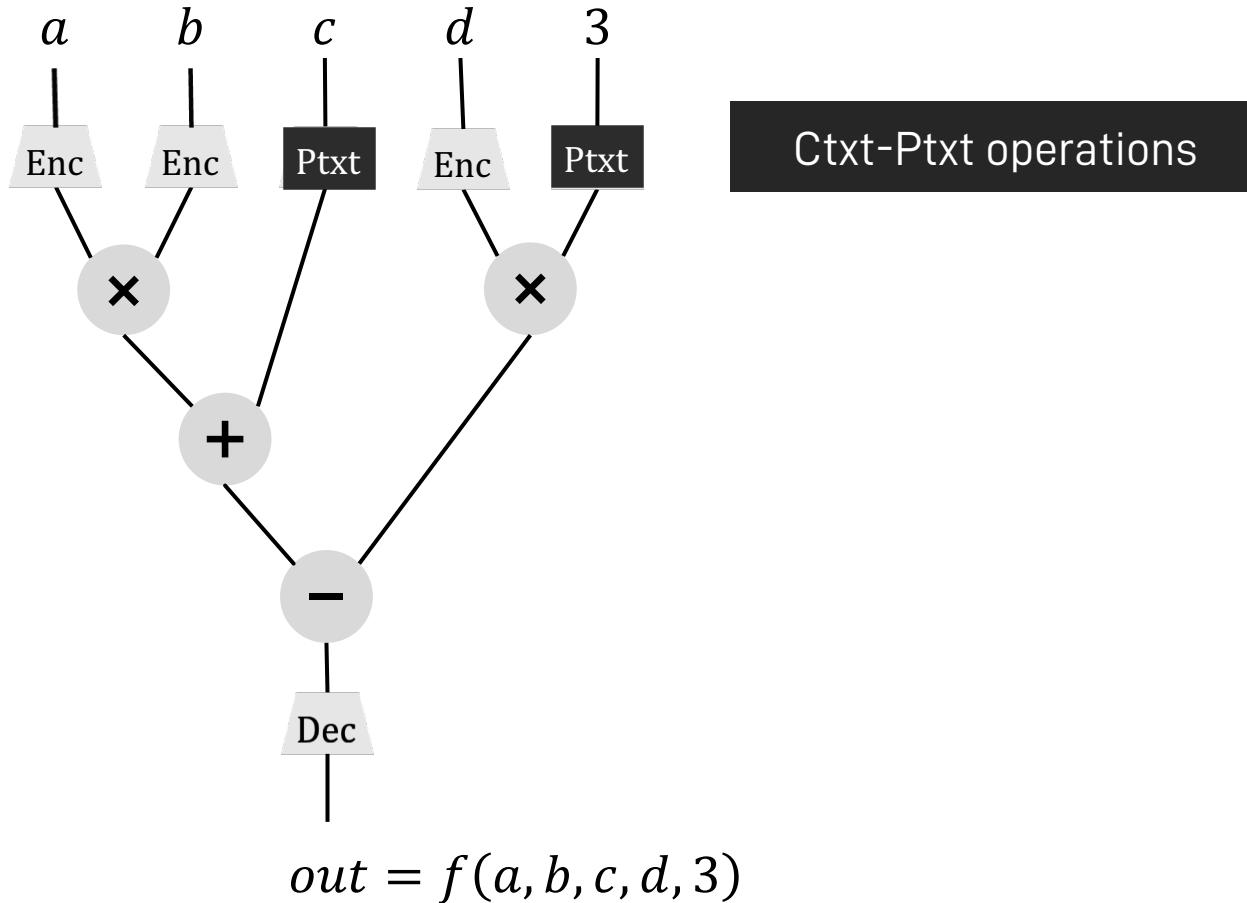


microsoft/SEAL  
Microsoft SEAL is an easy-to-use and powerful homomorphic encryption library.



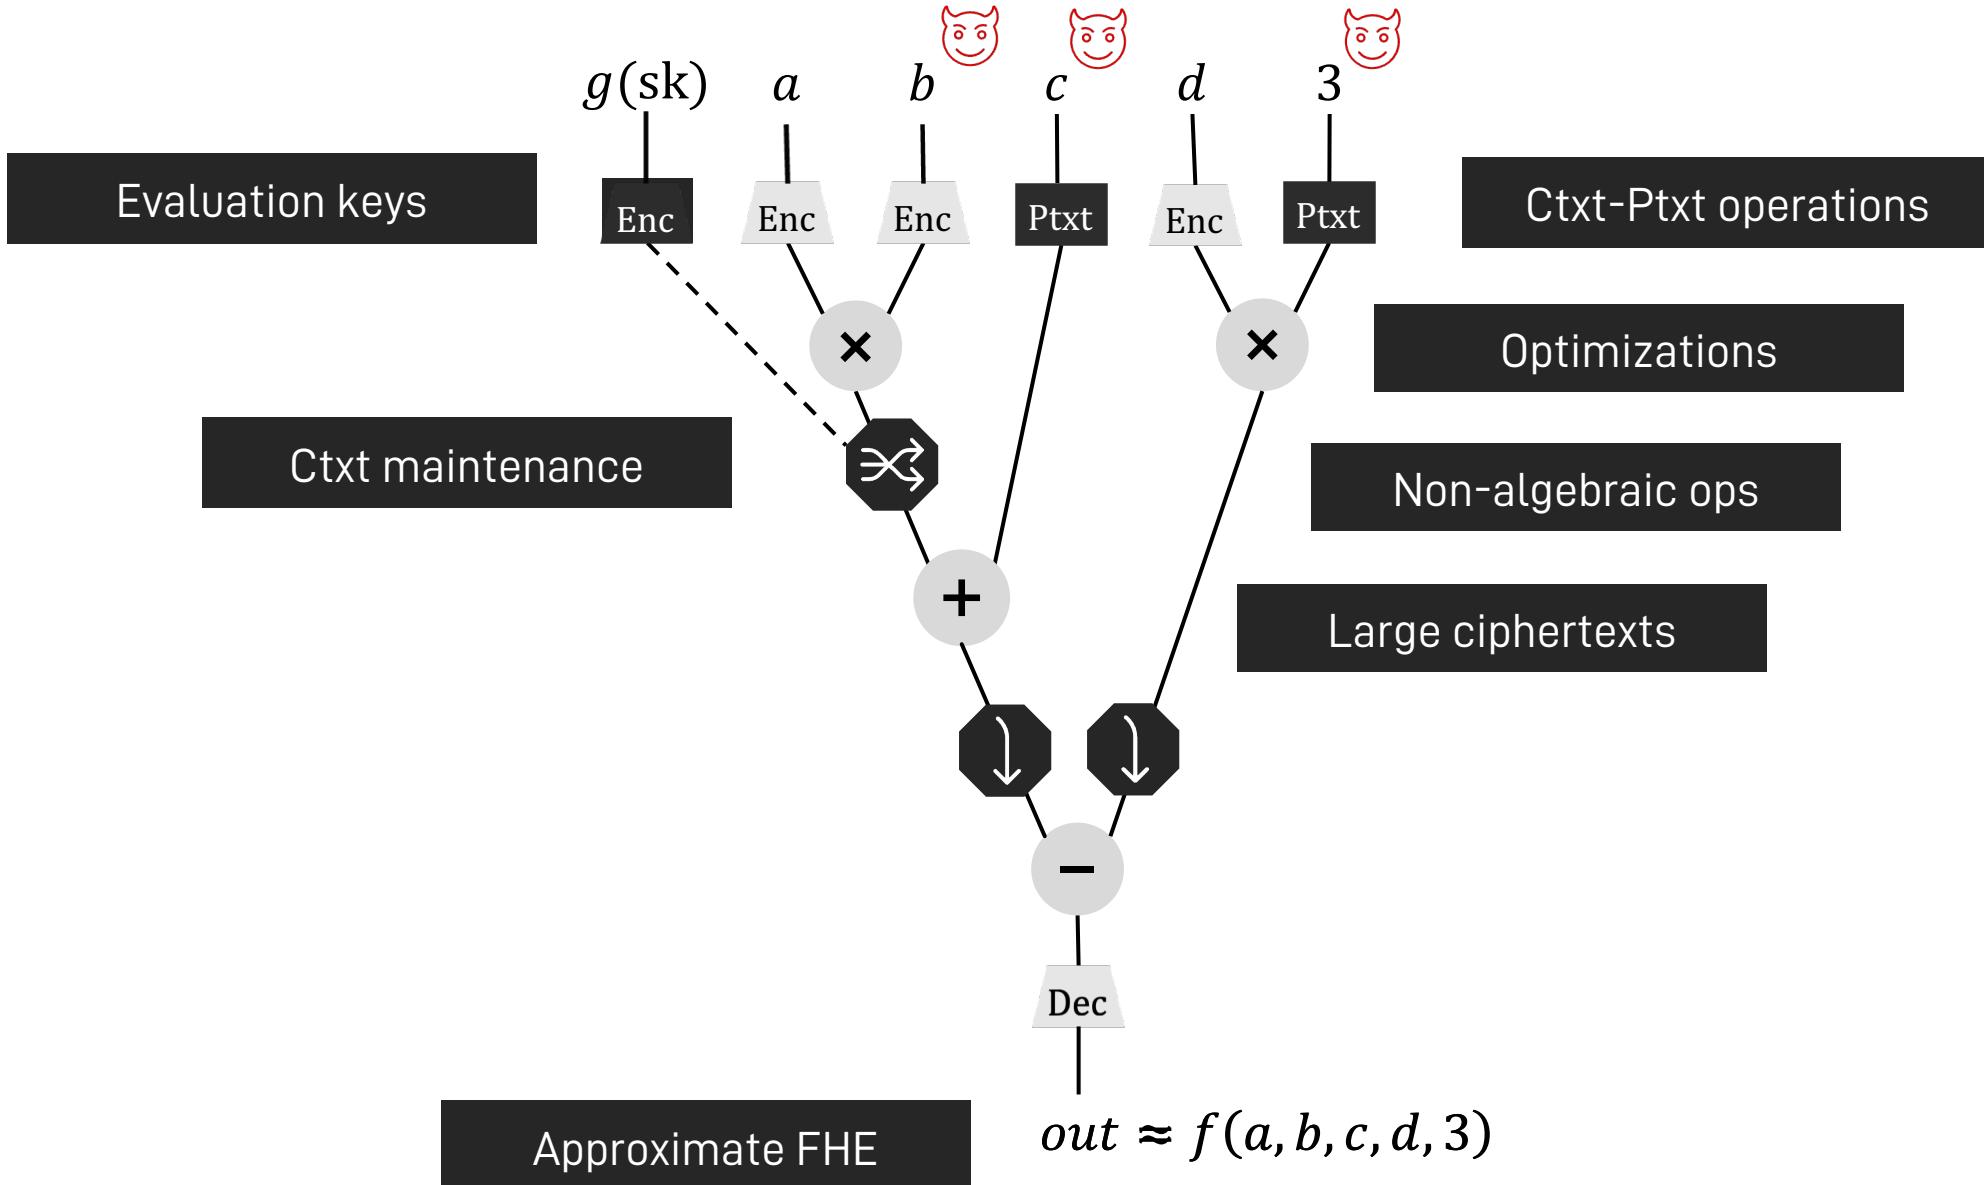
ZAMA  
TFHE-rs

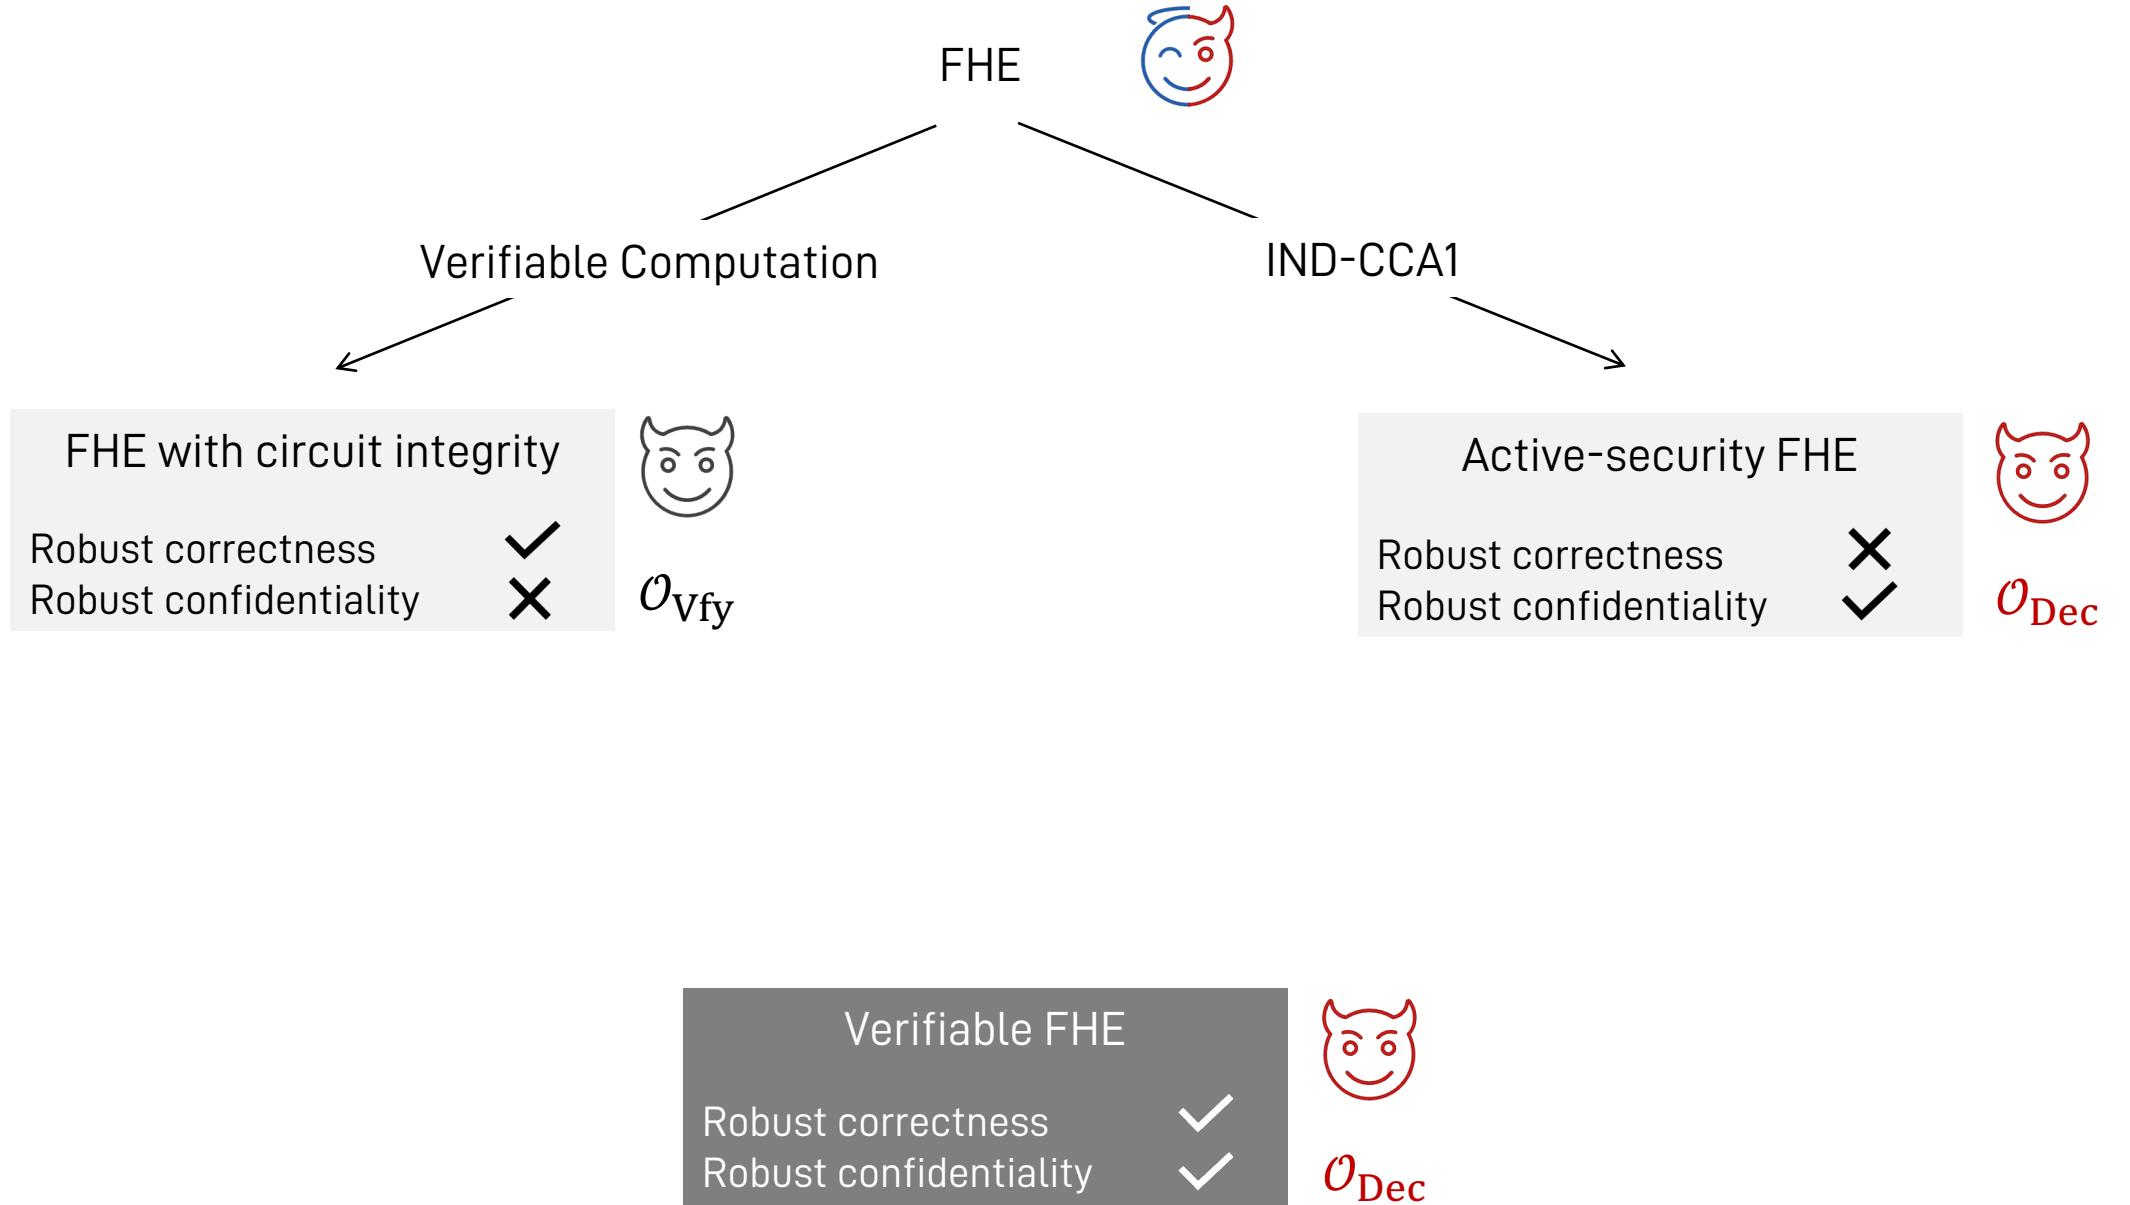
# Robustness challenges for modern FHE

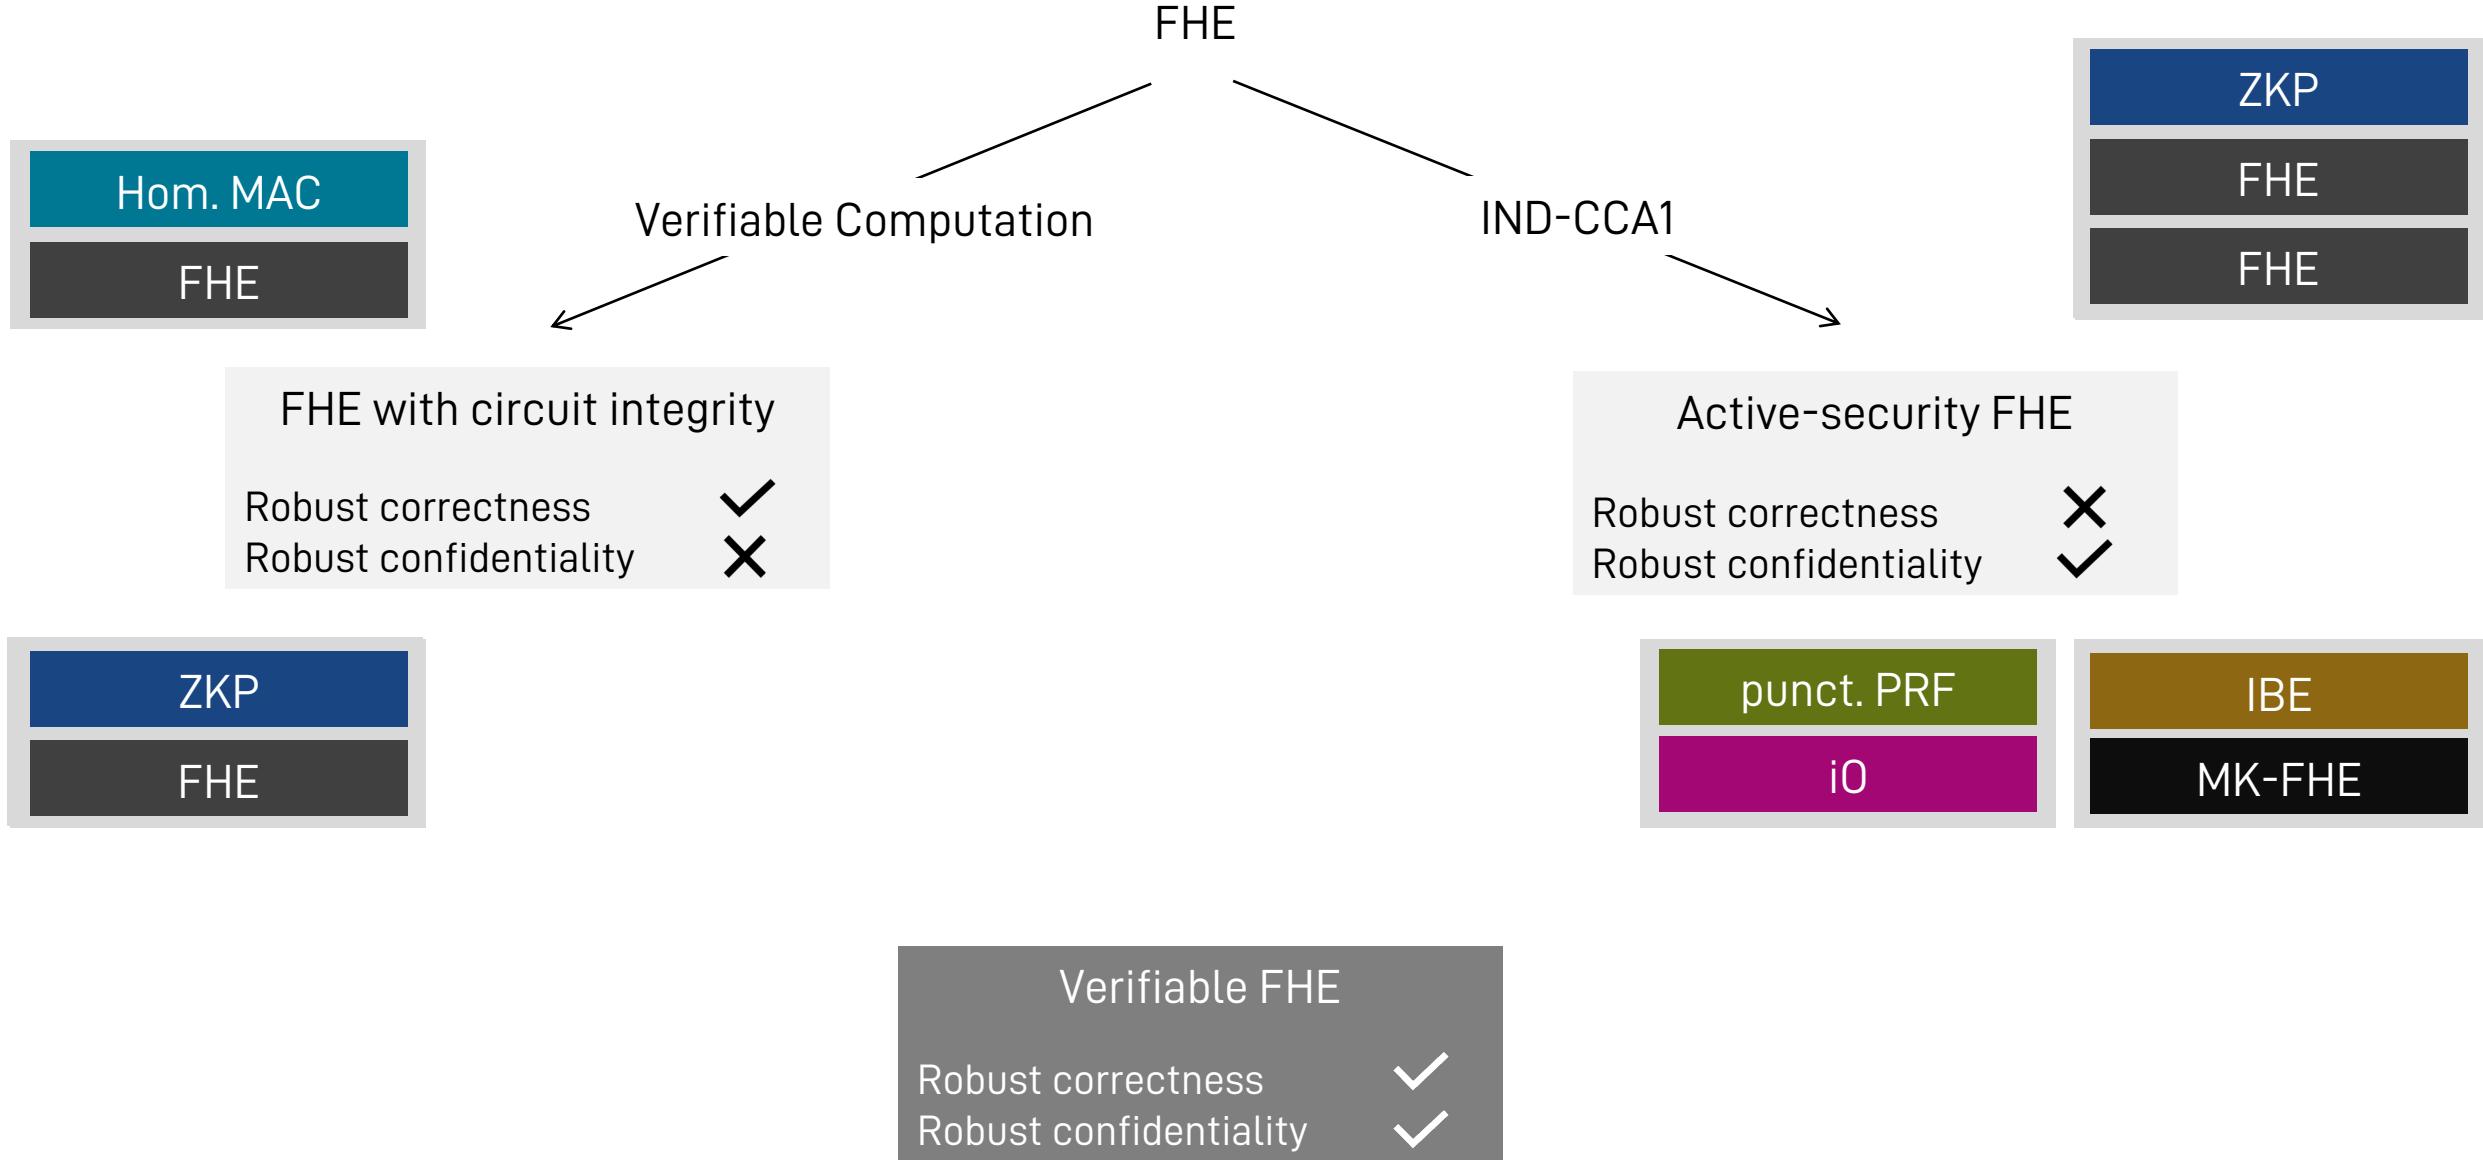


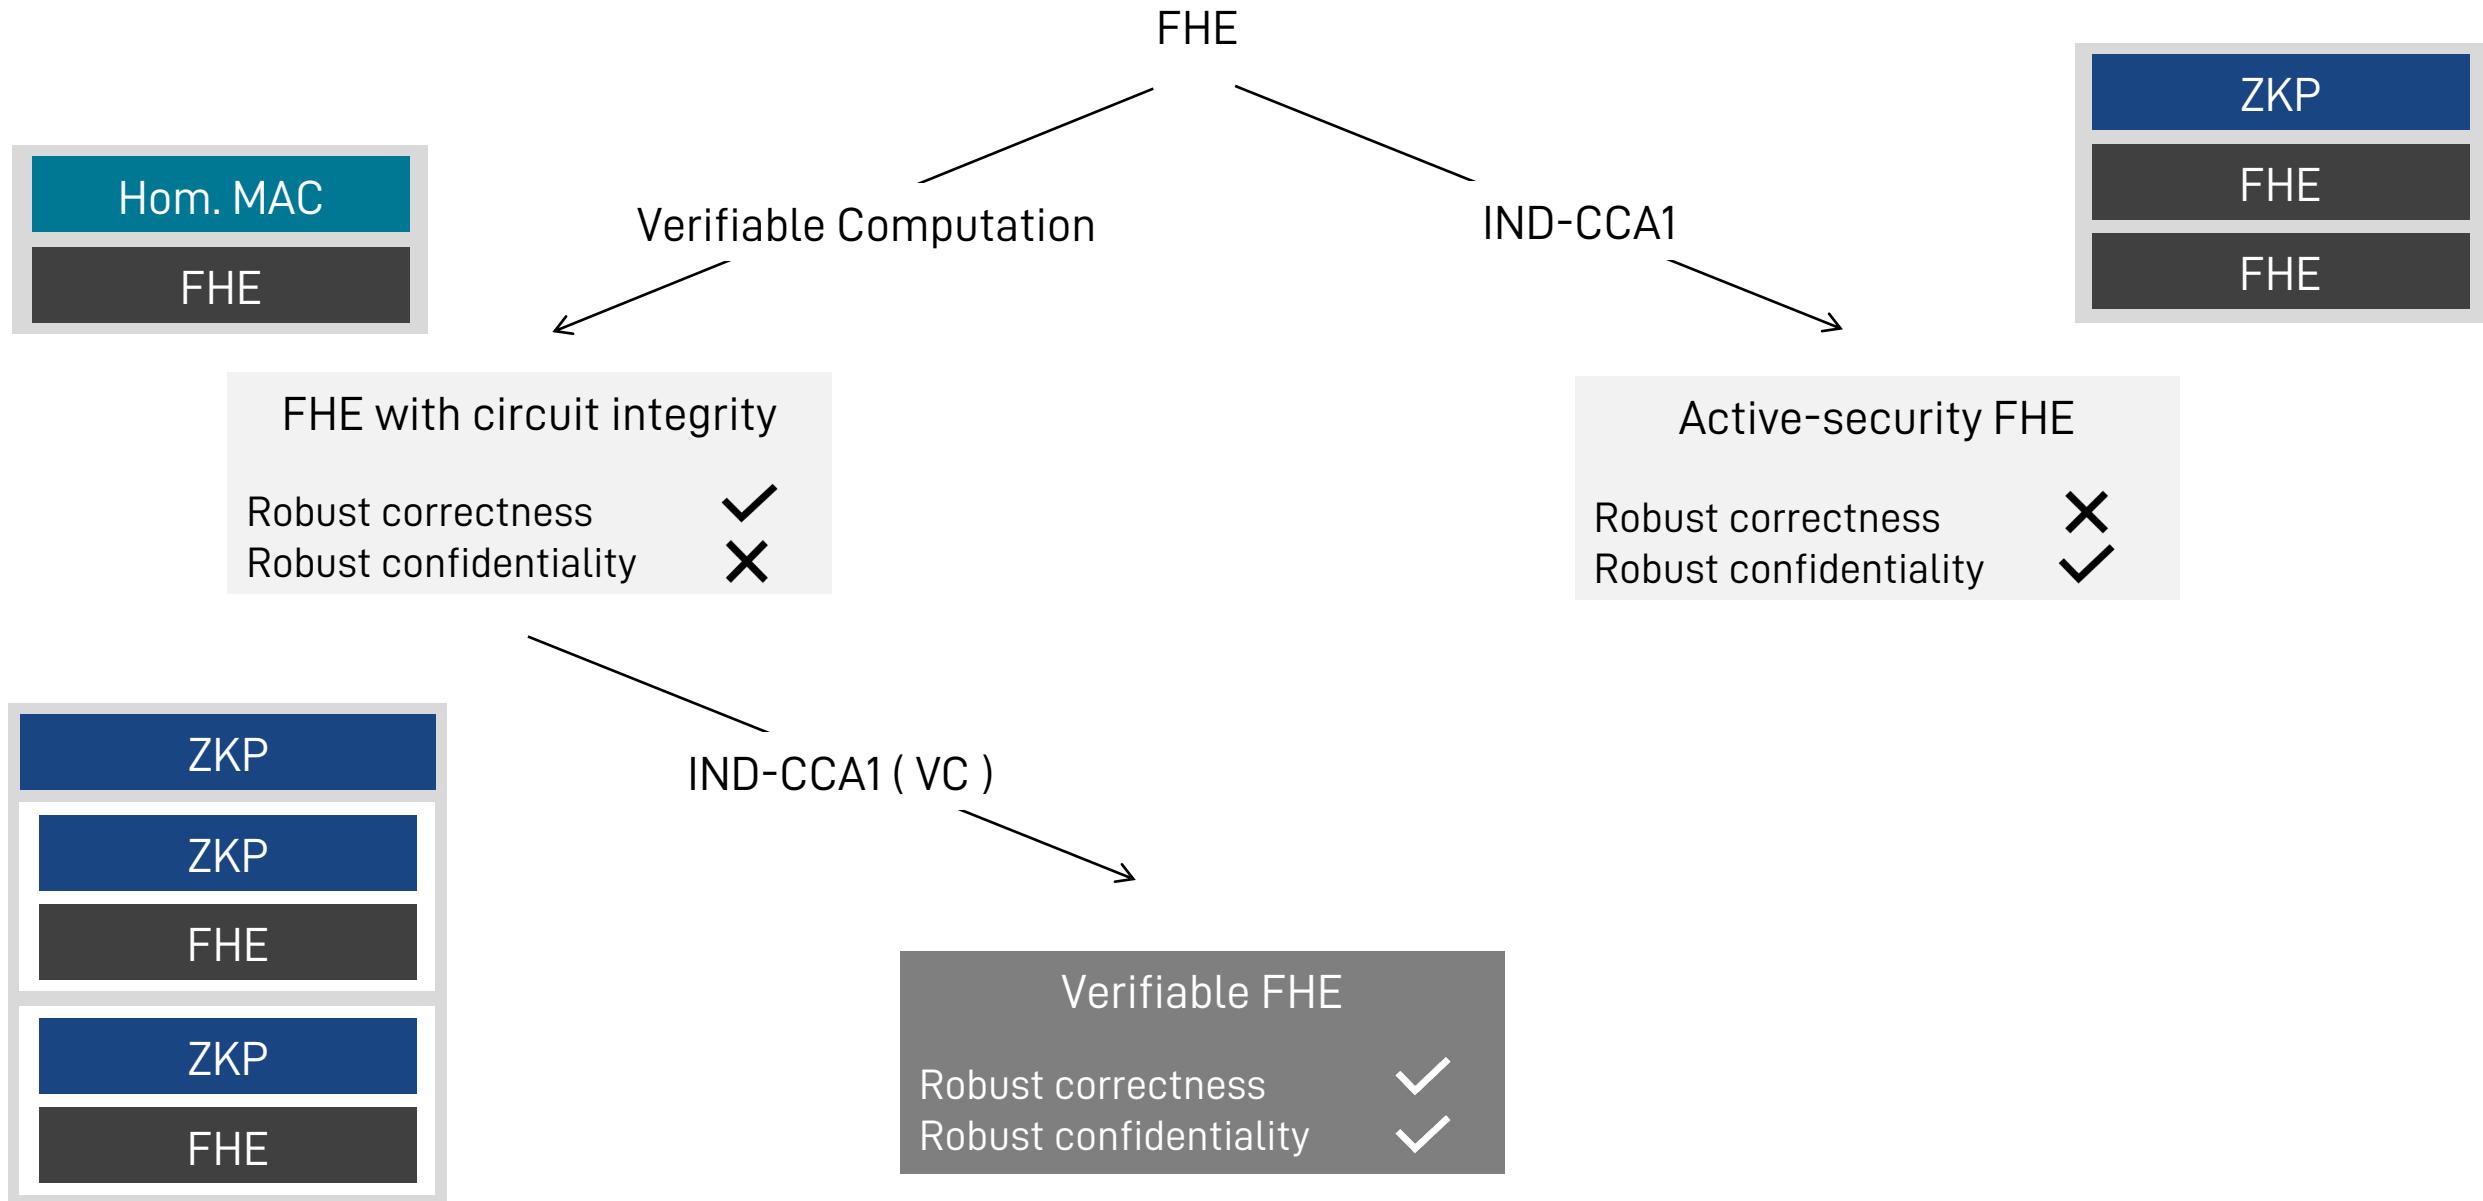
# Robustness challenges for modern FHE

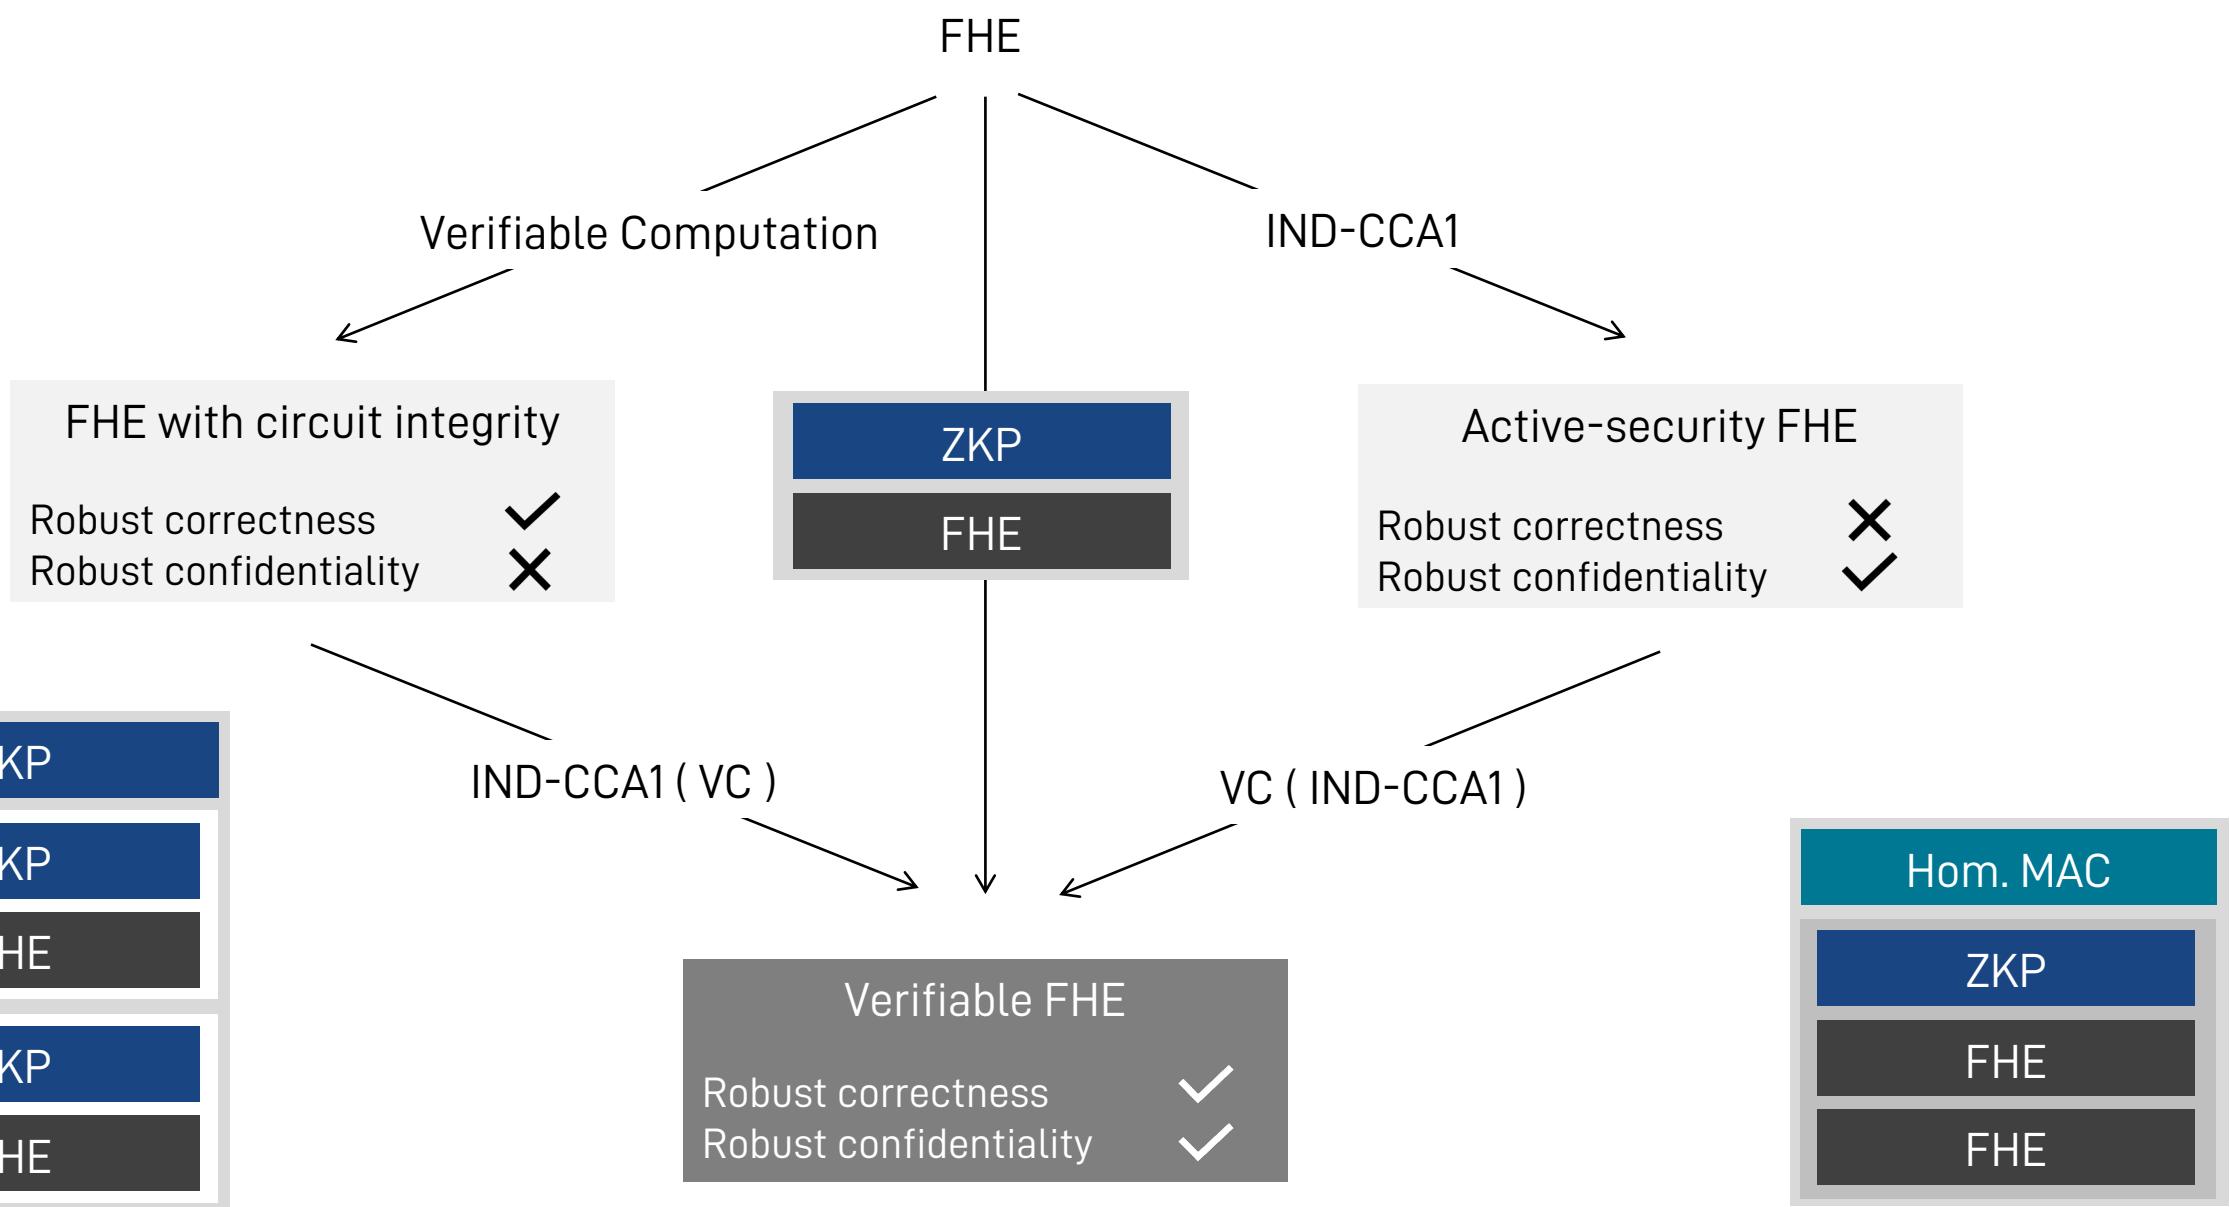
Server Inputs











## Starting from IND-CPA<sup>D</sup>

IND-CPA FHE

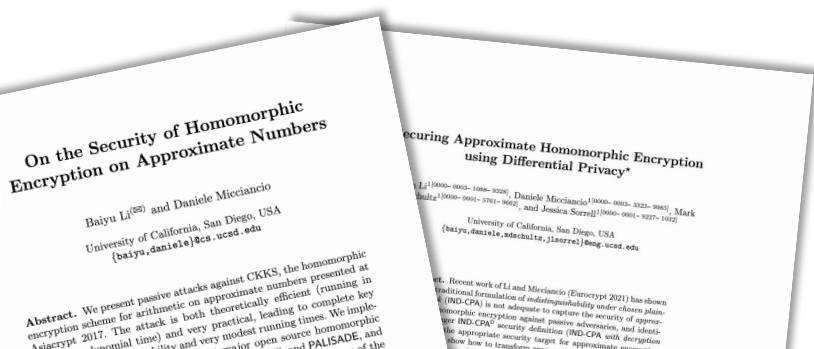
Exact FHE  
or +Diff. Priv.

IND-CPA<sup>D</sup> FHE

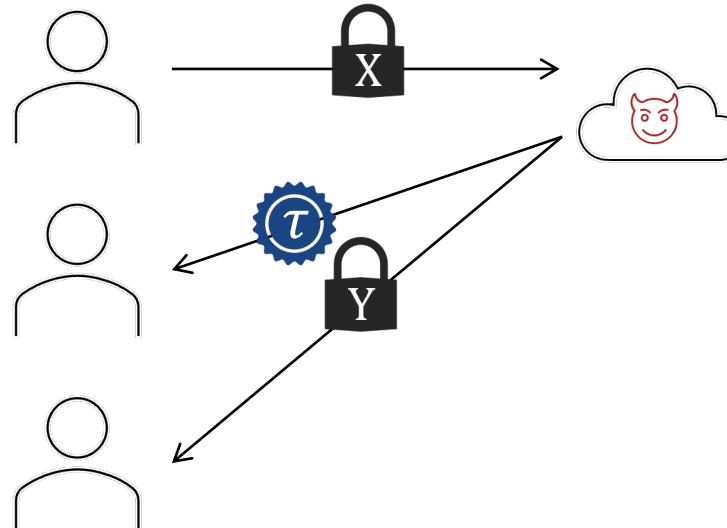
ZKP

FHE

Verifiable FHE



## Simplified setup



## Starting from IND-CPA<sup>D</sup>

IND-CPA FHE

Exact FHE  
or +Diff. Priv.

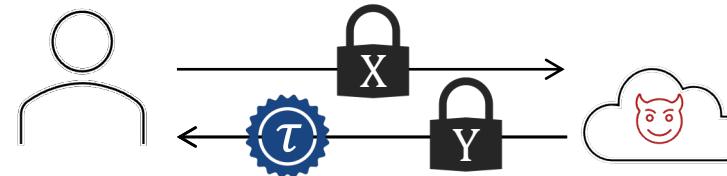
IND-CPA<sup>D</sup> FHE

ZKP

FHE

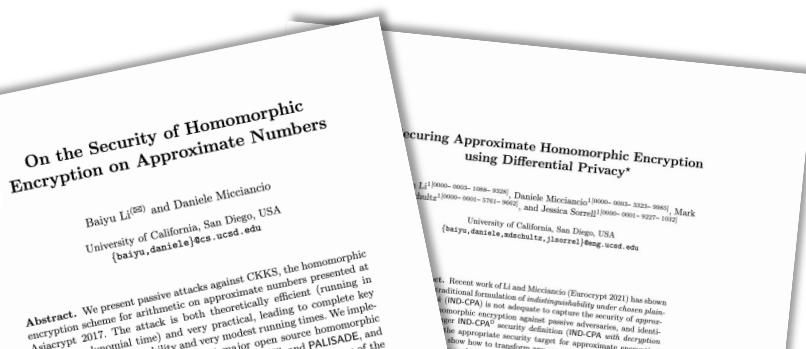
Verifiable FHE

## Simplified setup



## Plaintext server inputs only

$$\tau, \text{lock } Y = f(\text{lock } X, \text{lock } W)$$



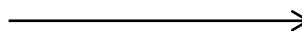
# Generic vFHE construction

KGen

Generate FHE keys ( $pk, sk$ )  
Generate keys for ZK proof system

Enc

$x$

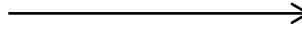


$c = \text{Enc}_{pk}(x)$

Eval

$c_1$

$w$



$c = \text{Eval}_{ek}(F, c_1, \dots, c_k, w)$

$\tau = \text{Prove}(c = \text{Eval}_{ek}(F, c_1, \dots, c_k, w))$   
 $\wedge w \text{ is valid plaintext}$

Verify  
Decrypt

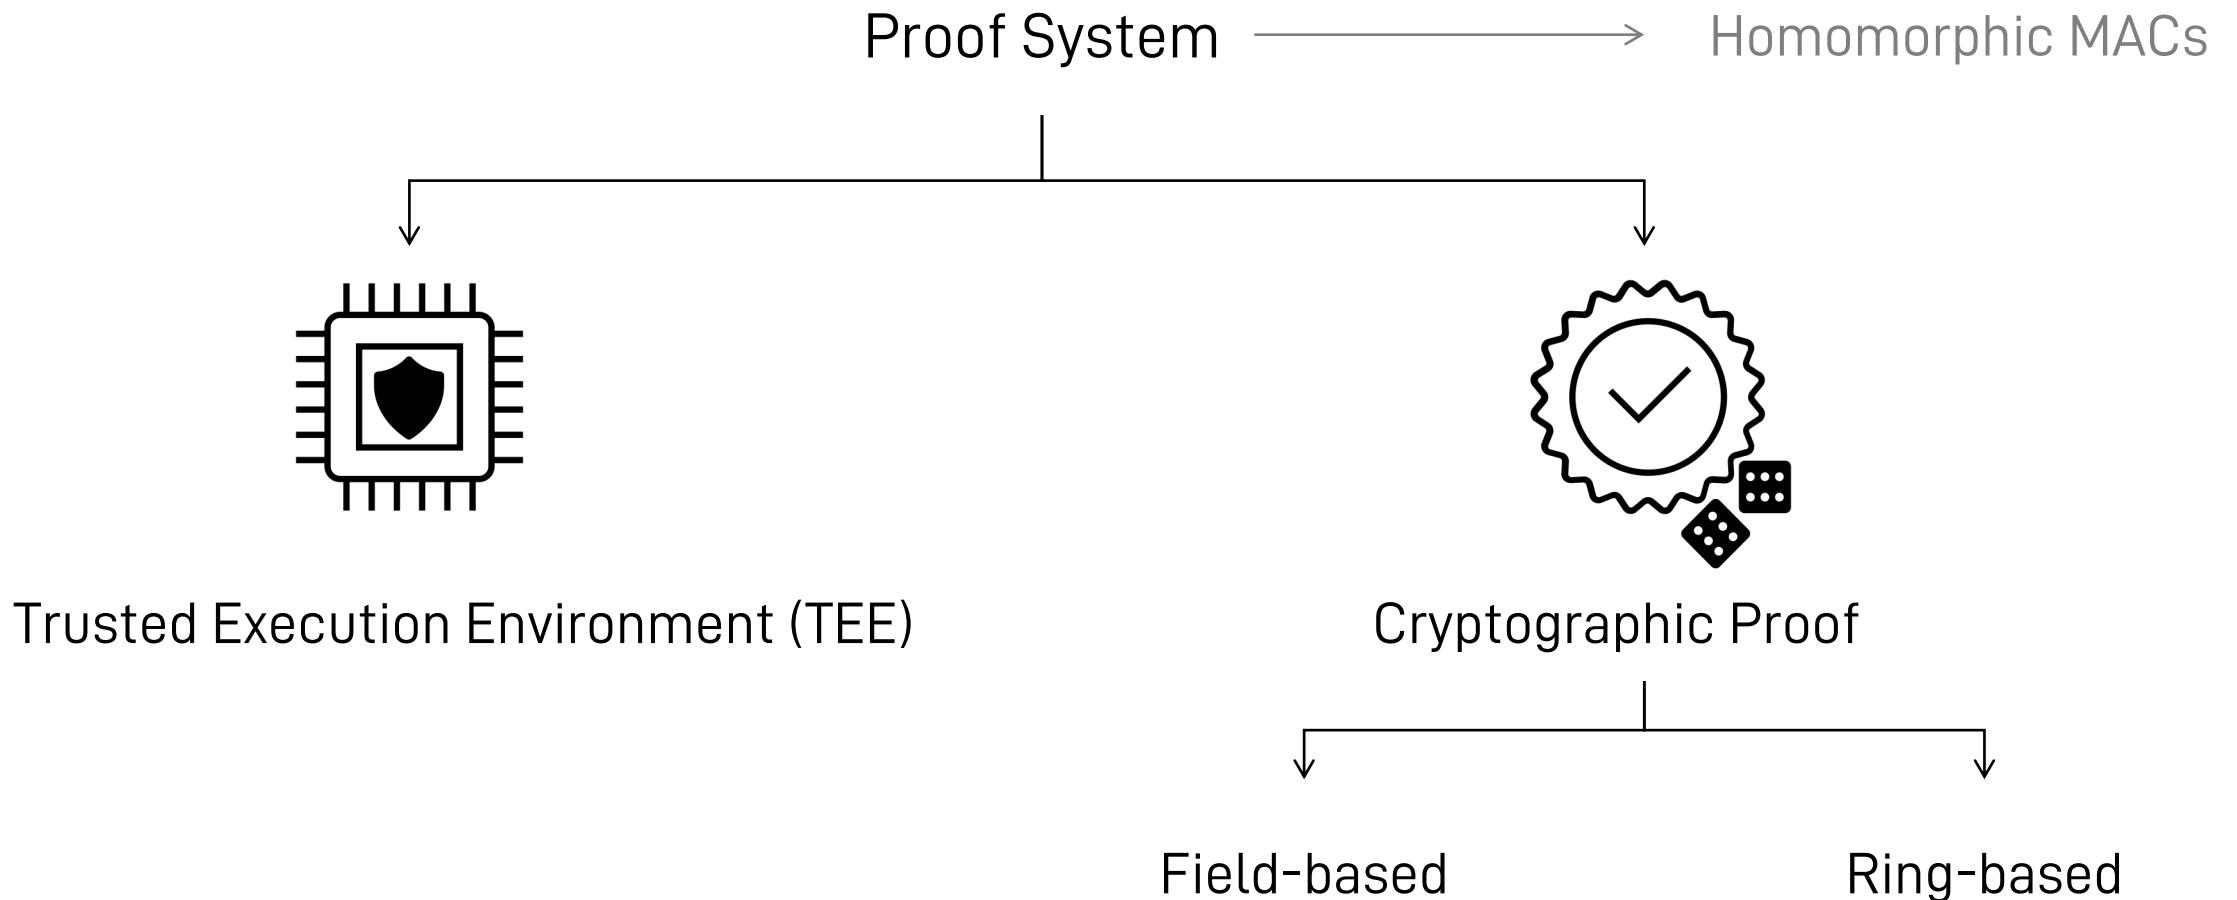
$c$

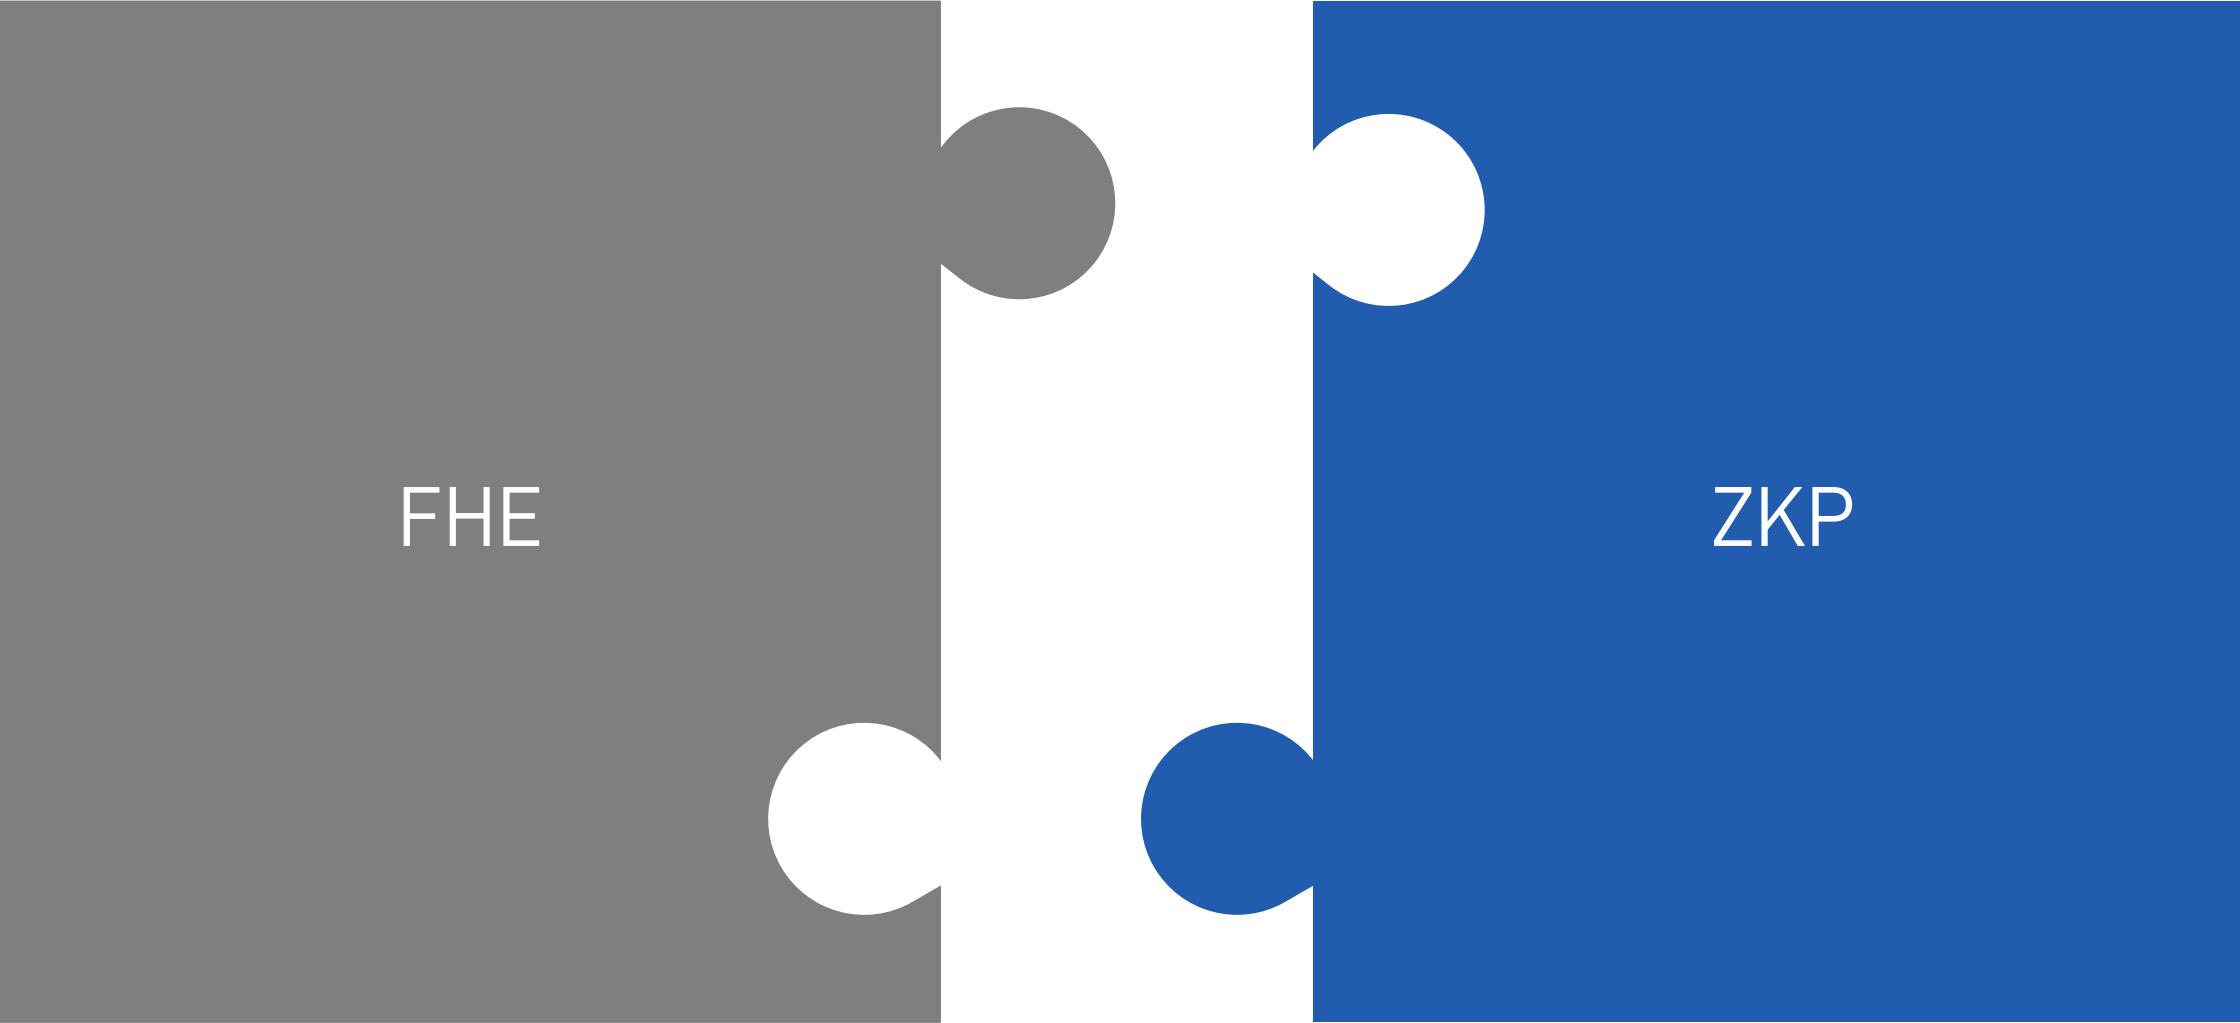
$\tau$

$\text{Verify}(\tau, c) = \text{true}$



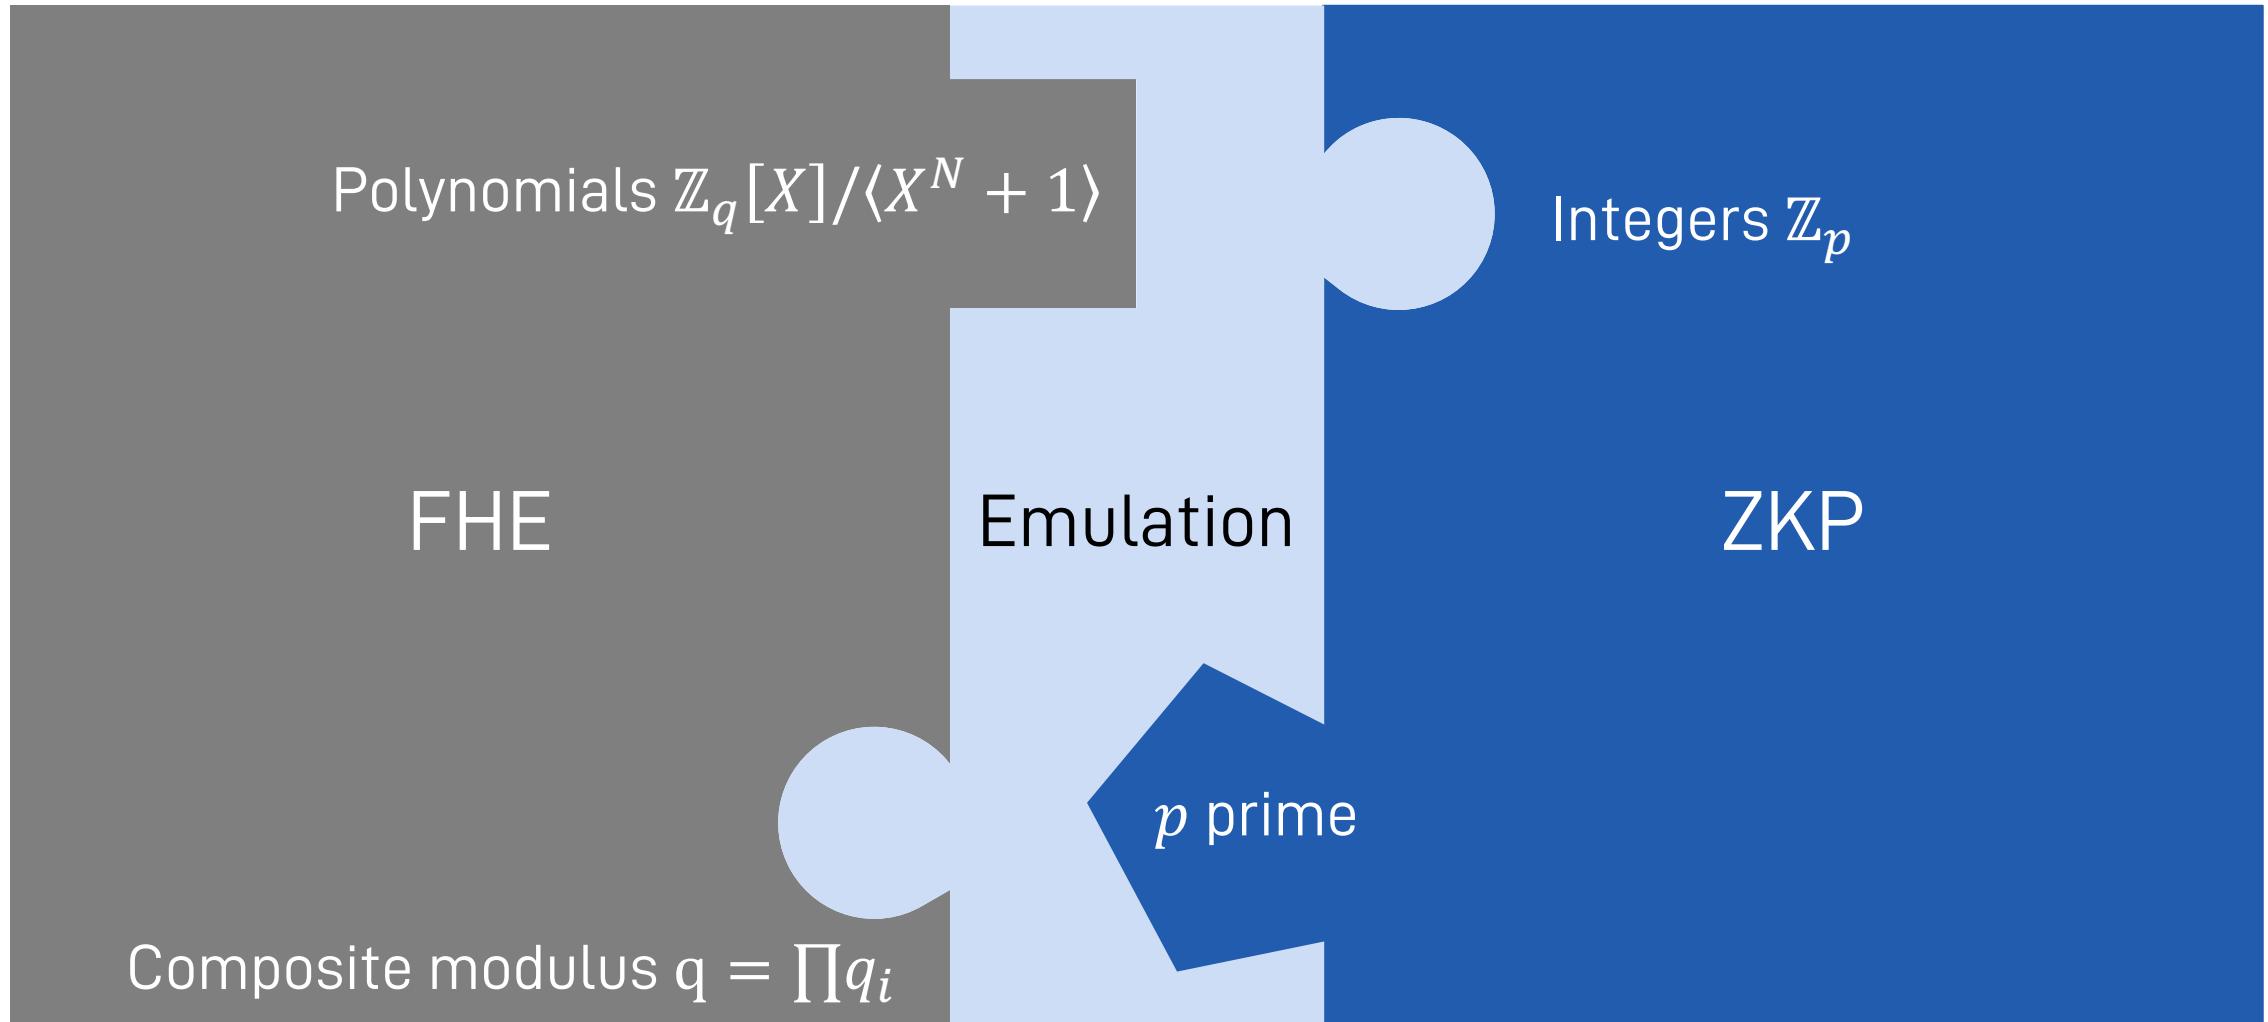
$y = \text{Dec}_{sk}(c)$

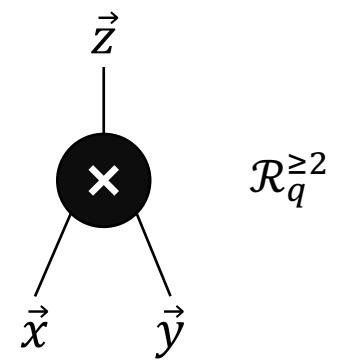


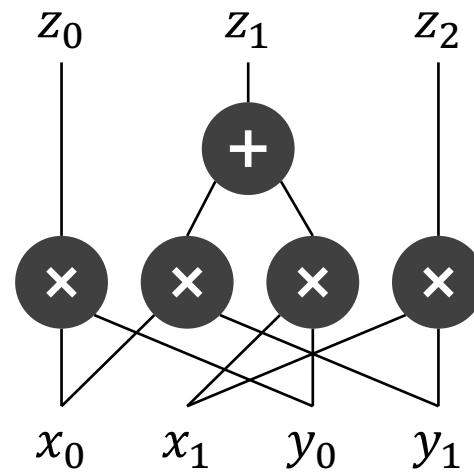


FHE

ZKP

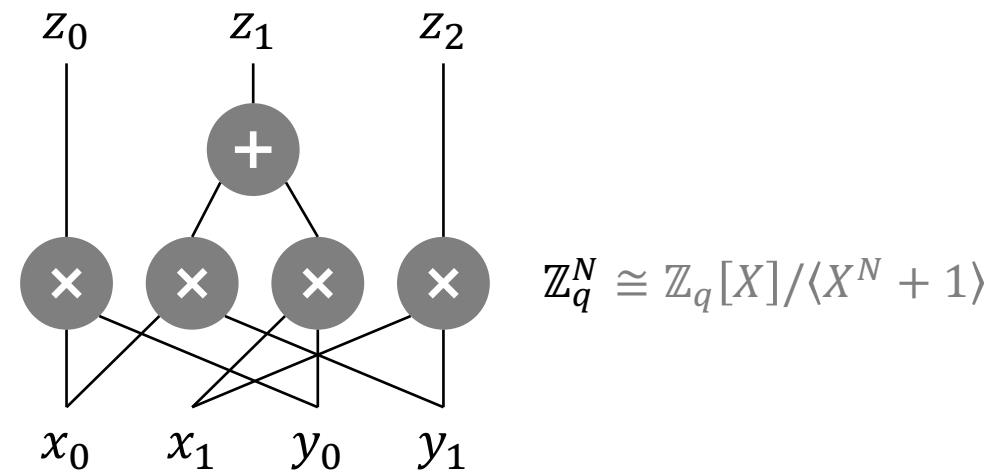


$\mathcal{R}_q^{\geq 2}$  $\mathbb{F}$ 

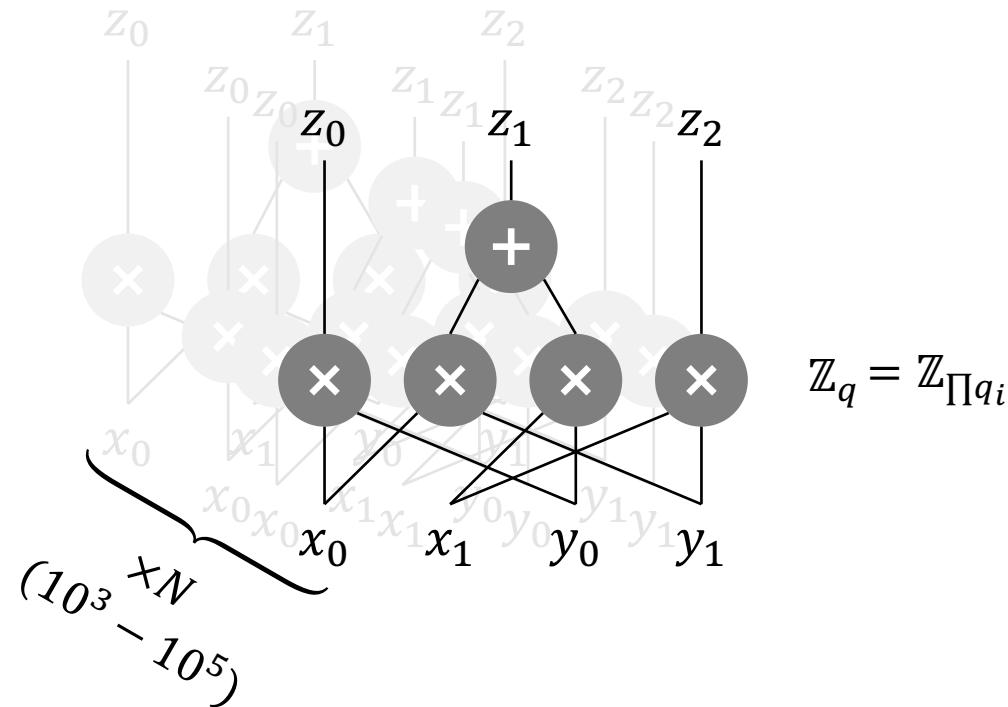
$\mathcal{R}_q^{\geq 2}$  $\mathcal{R}_q$  $\mathbb{F}$ 

$$\mathcal{R}_q = \mathbb{Z}_q[X]/\langle X^N + 1 \rangle$$

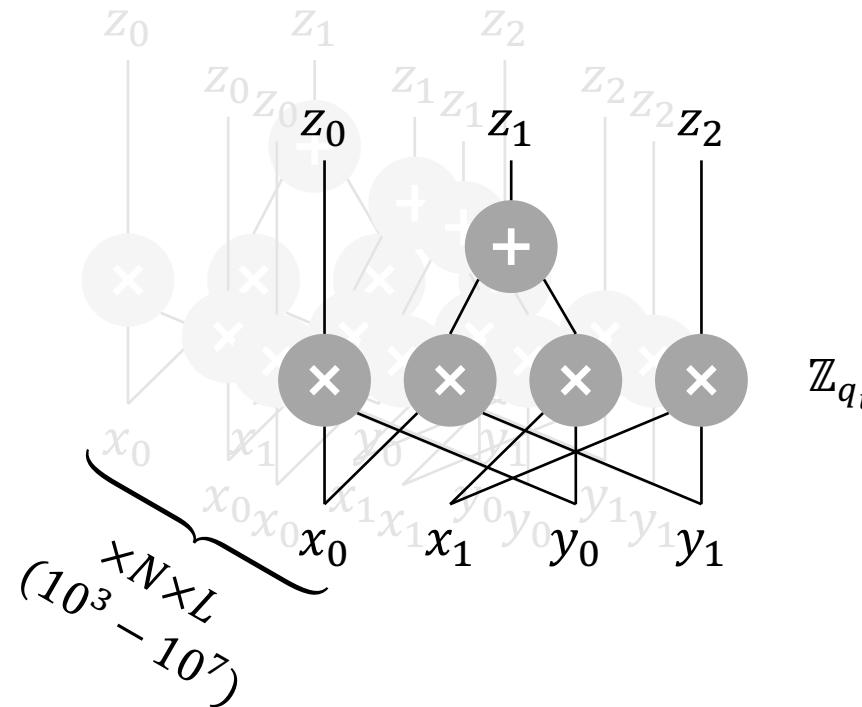
Overhead  $\times 4$

$\mathcal{R}_q^{\geq 2}$  $\mathcal{R}_q$  $\mathbb{Z}_q(\times N)$  $\mathbb{F}$ 

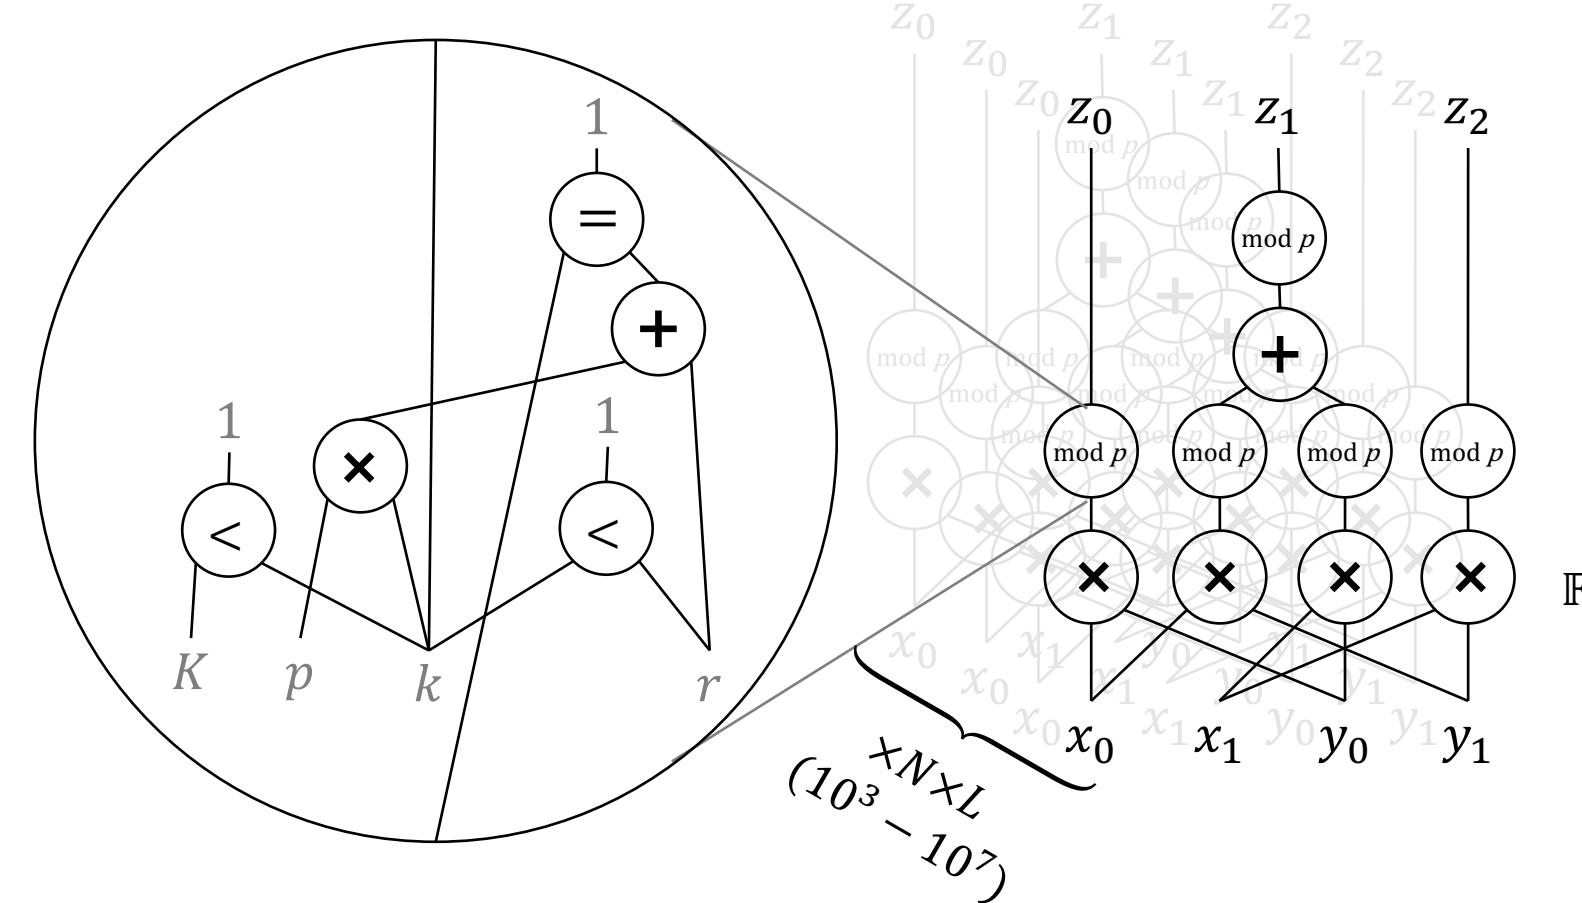
Overhead  $\times 4$

$\mathcal{R}_q^{\geq 2}$  $\mathcal{R}_q$  $\mathbb{Z}_q(\times N)$  $\mathbb{F}$ 

Overhead  $\times 32'000$

$\mathcal{R}_q^{\geq 2}$  $\mathcal{R}_q$  $\mathbb{Z}_q(\times N)$  $\mathbb{Z}_{q_i}(\times N \times L)$  $\mathbb{F}$ 

Overhead  $\times 130'000$

$\mathcal{R}_q^{\geq 2}$  $\mathcal{R}_q$  $\mathbb{Z}_q(\times N)$  $\mathbb{Z}_{q_i}(\times N \times L)$  $\mathbb{F}$ 

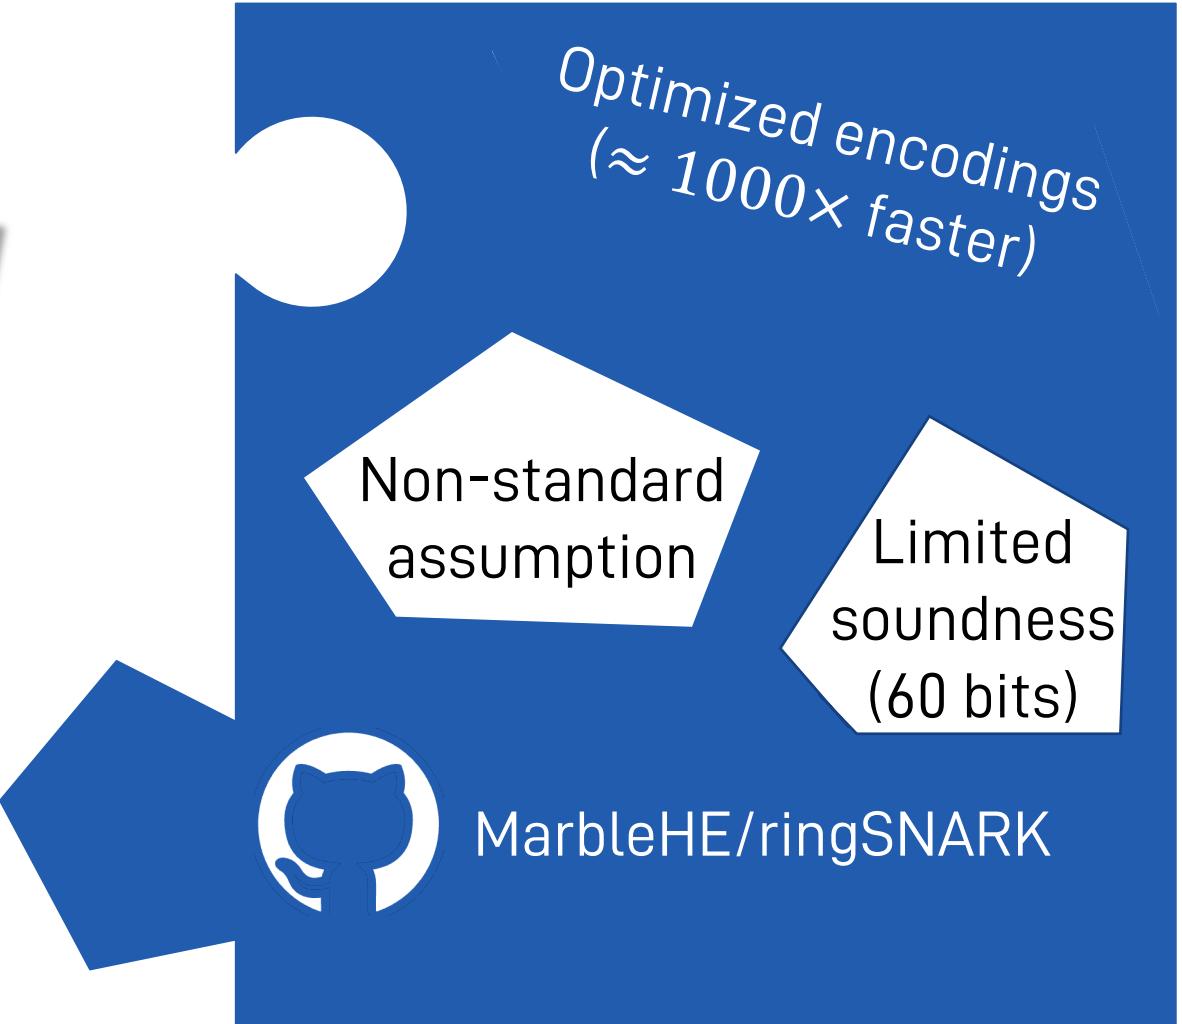
Overhead  $\times 16'000'000$

## A better fit: ring-based ZKPs

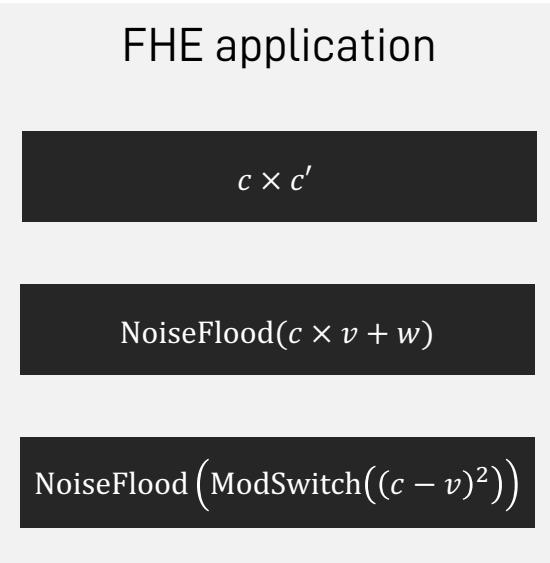




# State-of-the-art ring-ZKP: Rinocchio



# Evaluation



SEAL/app.cpp

```
evaluator.sub_plain_inplace(x_encrypted, w_plain);
evaluator.square_inplace(x_encrypted);
evaluator.mod_switch_to_next_inplace(x_encrypted);

for (int i = 0; i < SEC_PARAM; i++) { // Noise flooding
    evaluator.multiply_plain(zeros_encrypted[i], b_plain[i], tmp);
    evaluator.add_inplace(x_encrypted, tmp);
}
```

Rinocchio/app.cpp

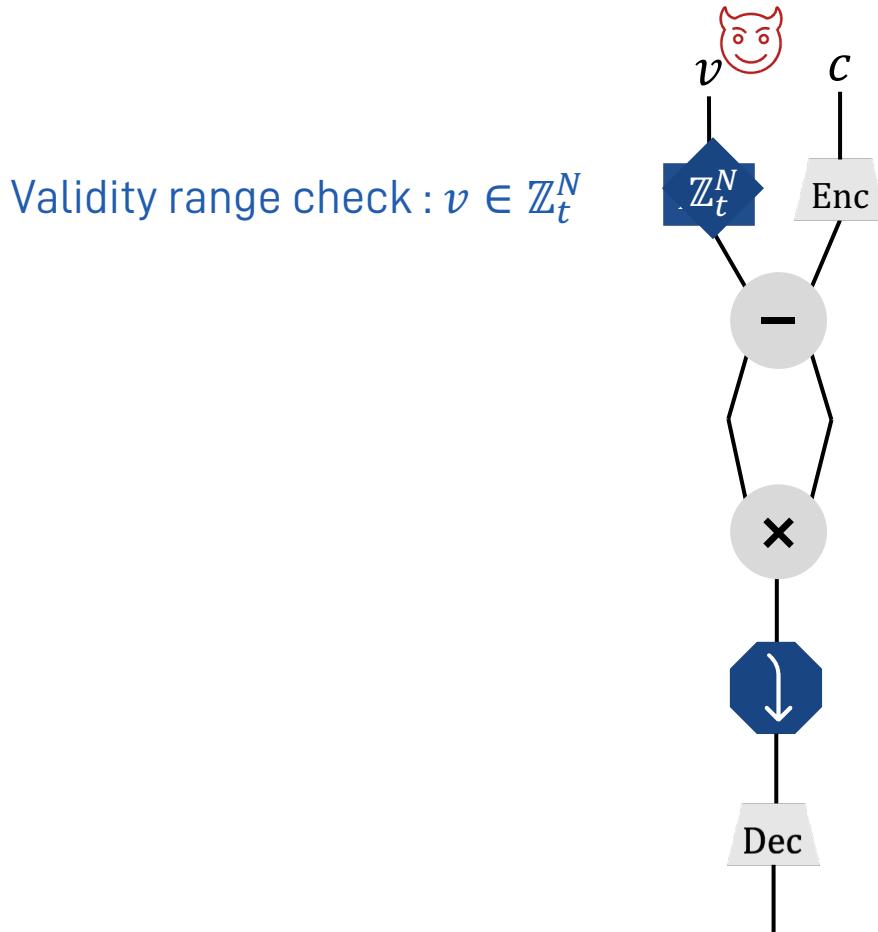
Bulletproofs/app.rs

Circom/app.circom

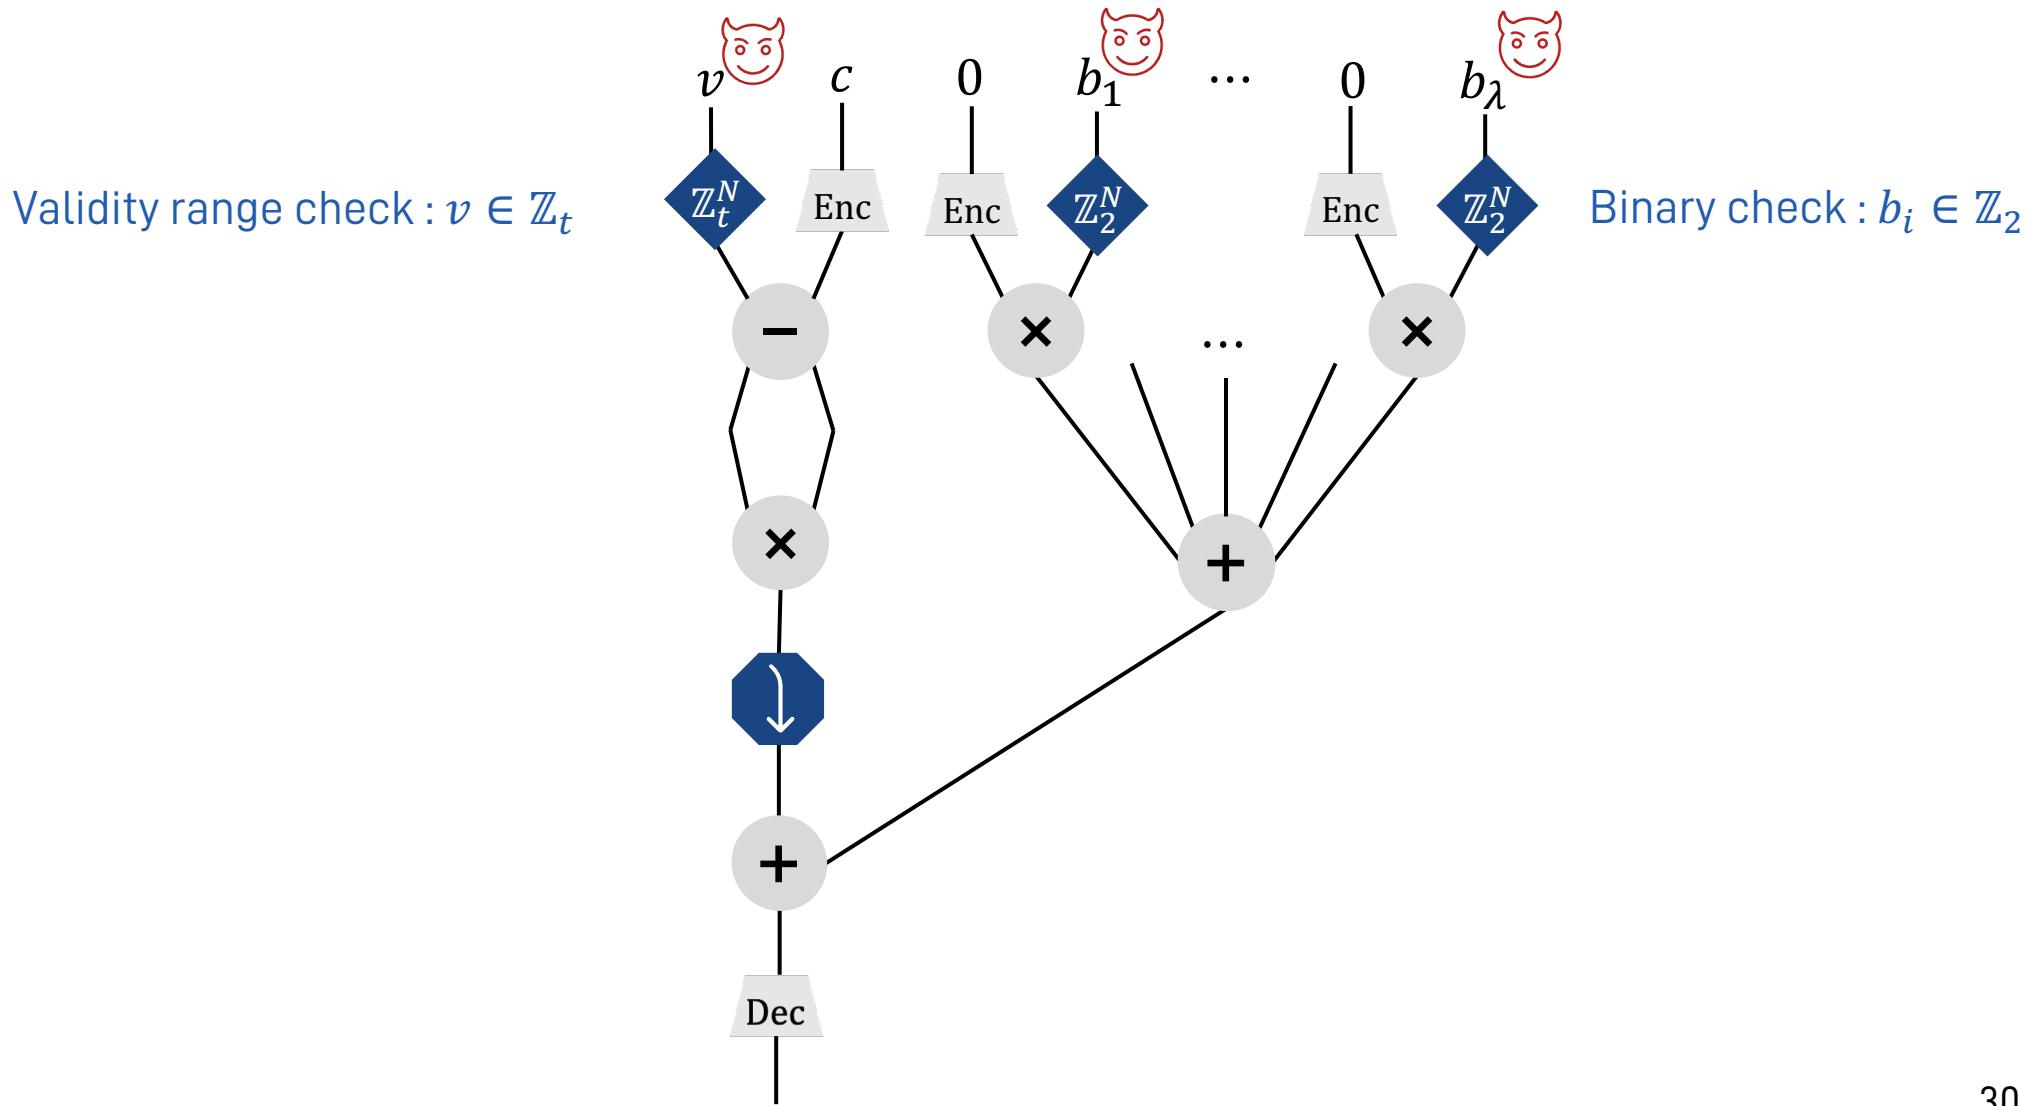
```
signal input in[2][l][n], w[n], noise[secpar][3][l_][n], b[secpar];
LtConstantN(t, n)(w); // Input checks
signal w_ntt[l][n] ← NTTsPlain(l, n, qs)(w);

signal diff[2][l][n] ← FastSubPlain(l, n, qs)(in, w_ntt);
signal mul[3][l][n] ← SquareCtxt(l, n, qs)(diff);
signal modsw[3][l_][n] ← ModSwitchToNext(l, n)(mul);
signal output out[3][l_][n]; ← NFlood(secpar, l_, n, qs)(modsw, noise, b);
```

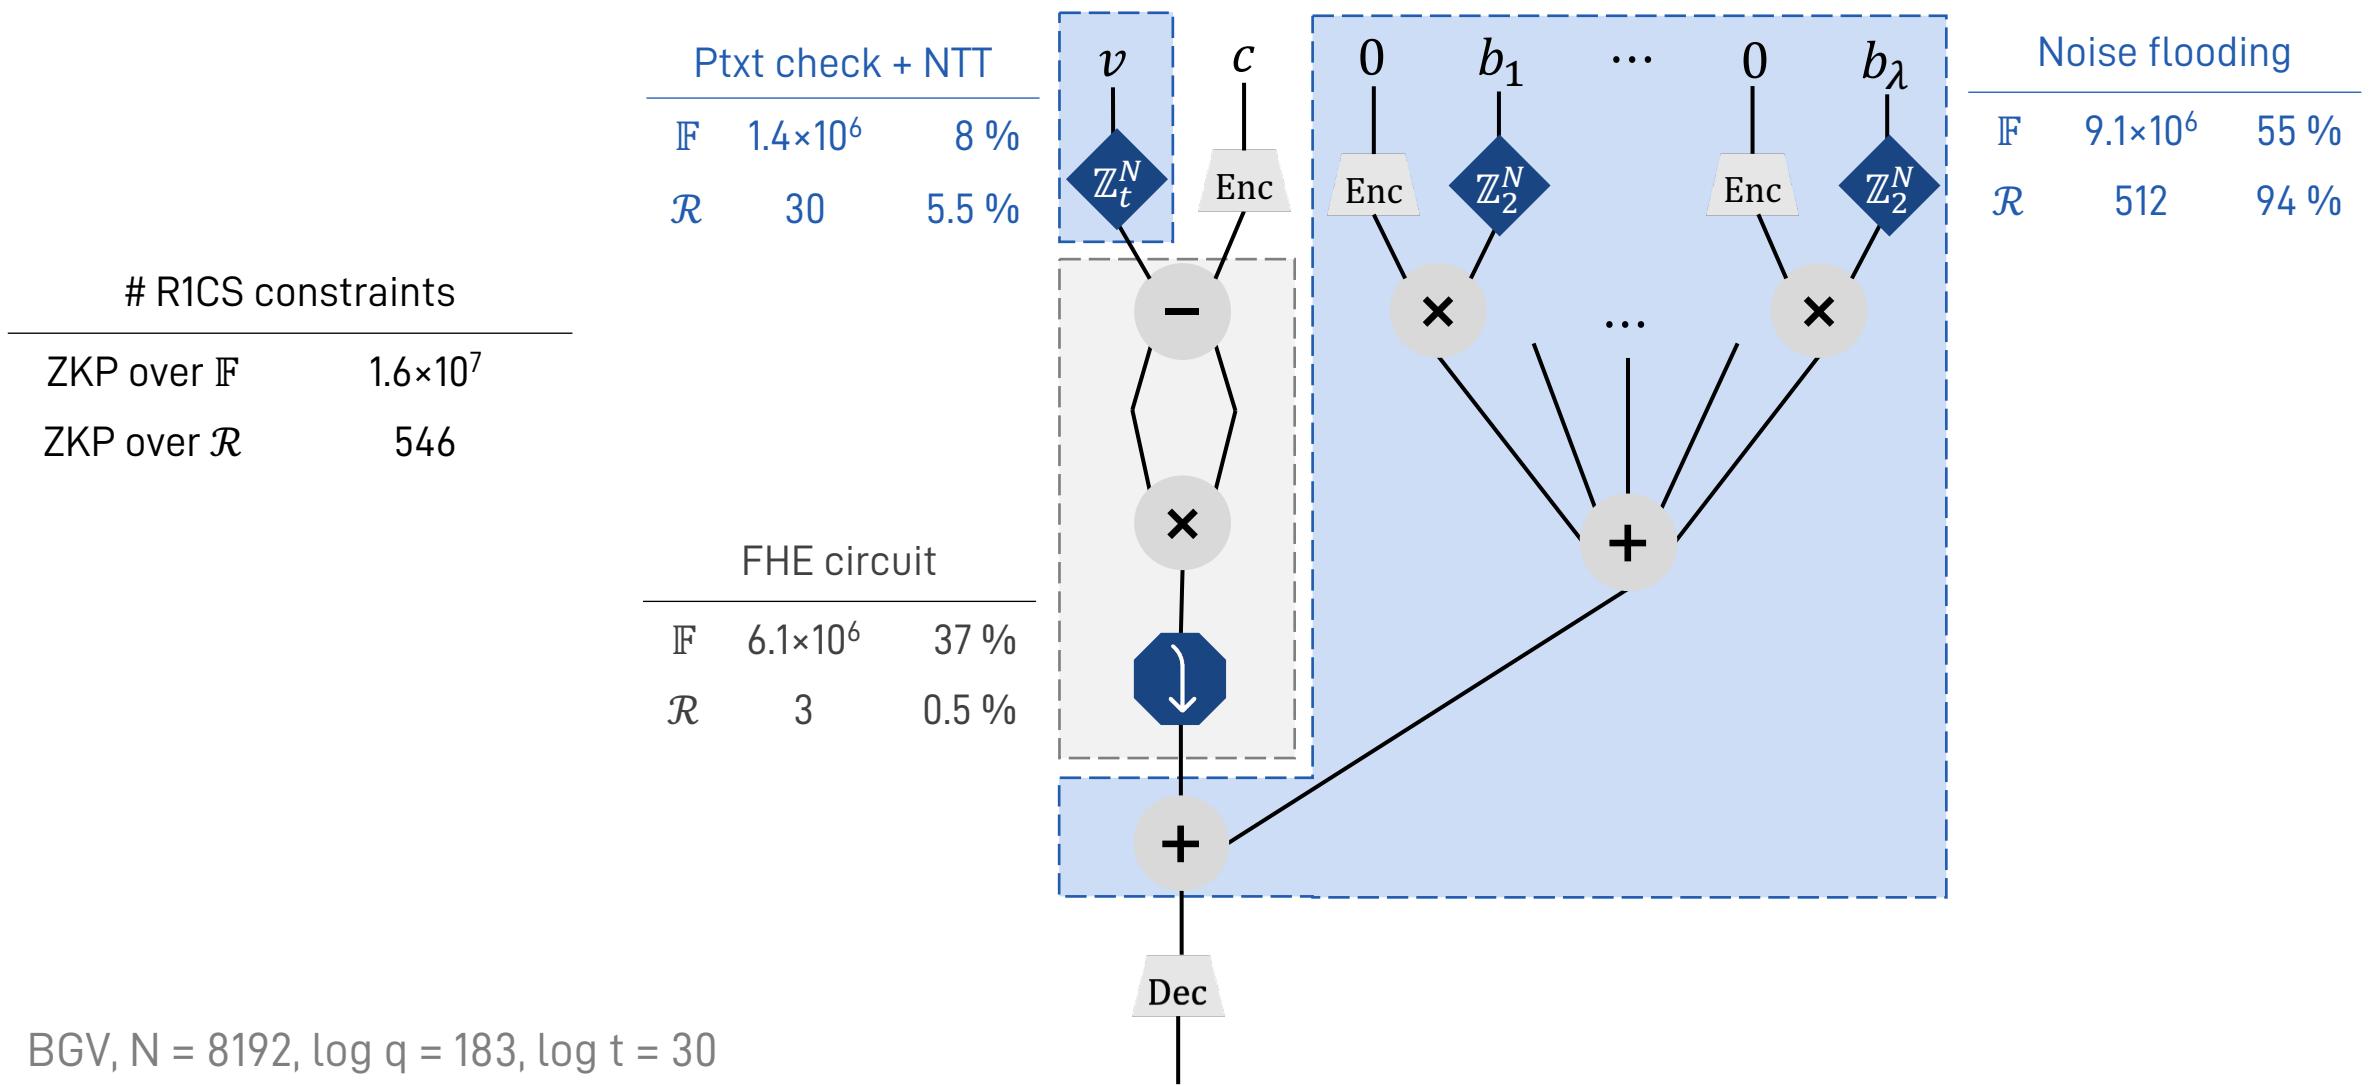
$$f = \text{ModSwitch}((c - v)^2)$$



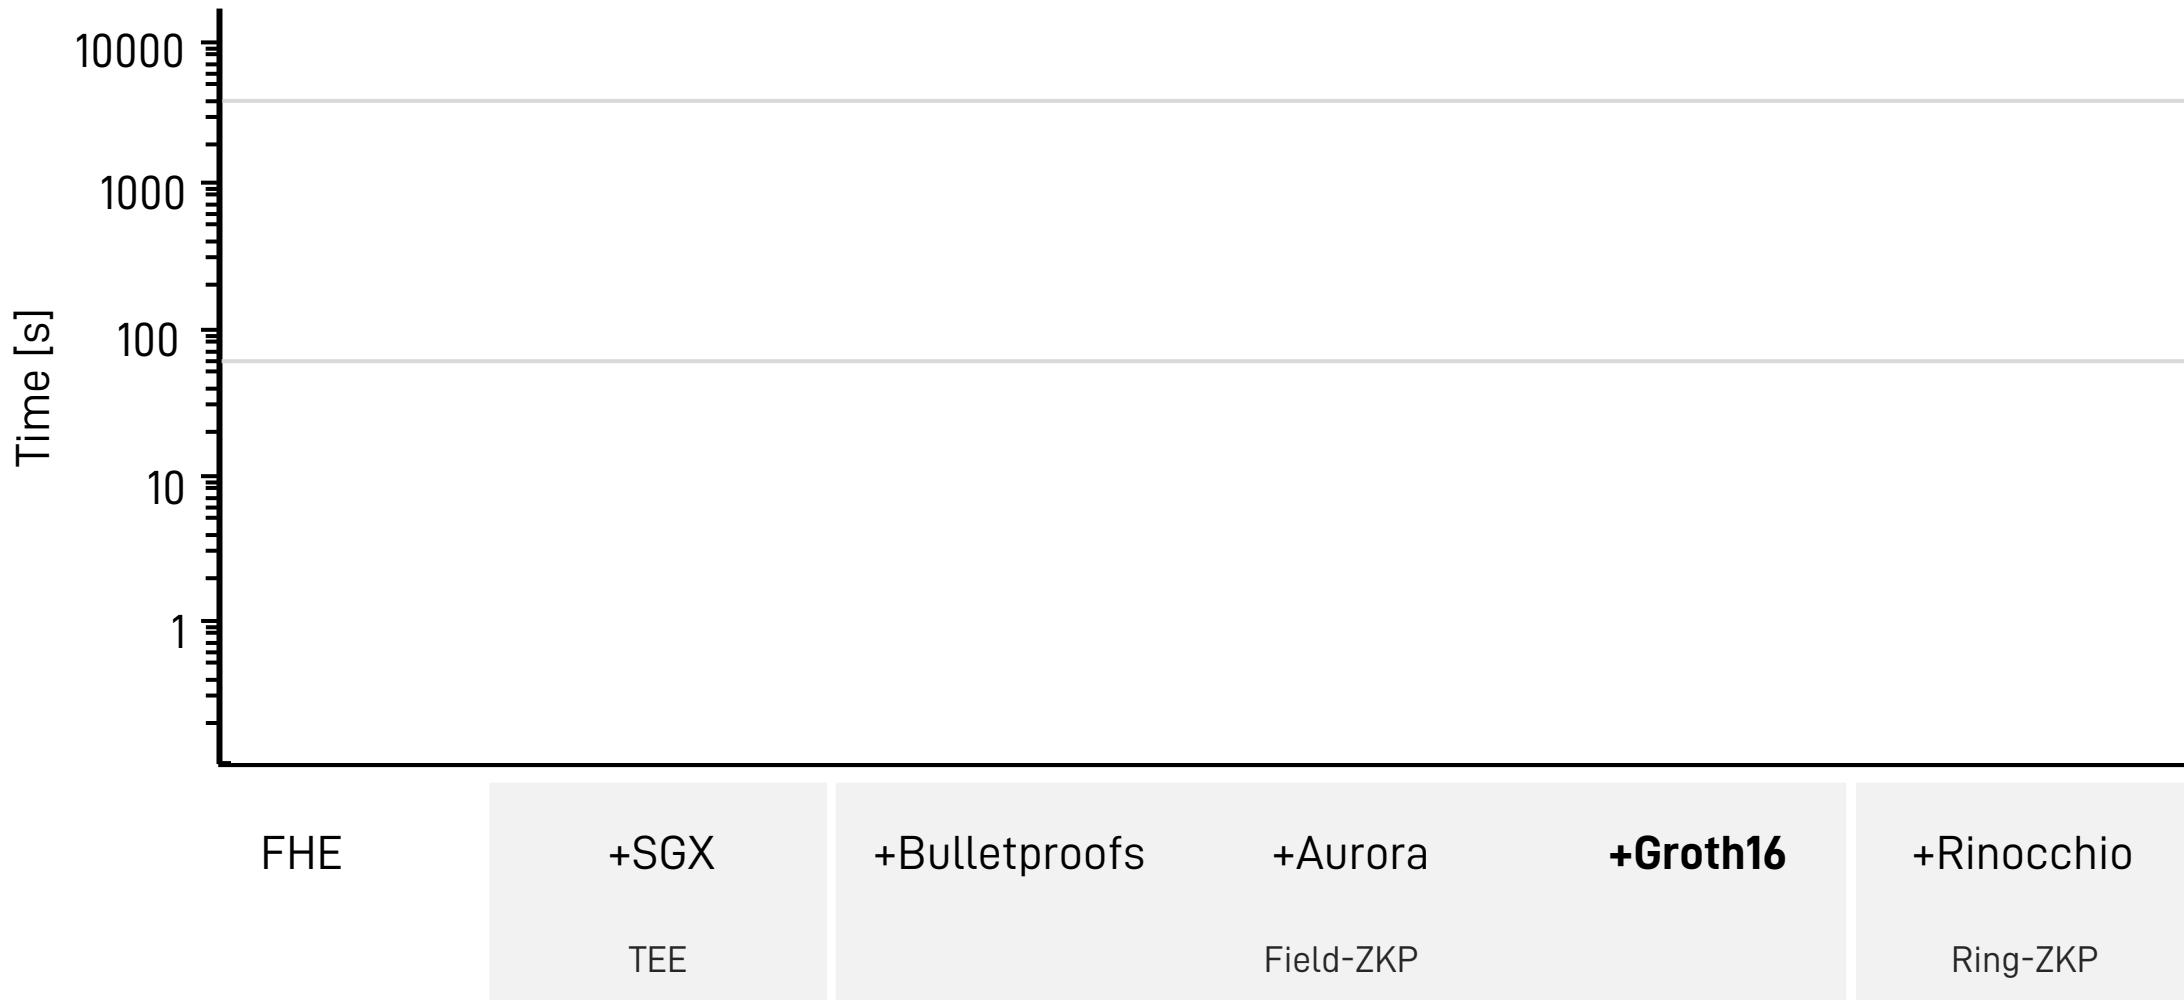
$$f = \text{NoiseFlood}\left(\text{ModSwitch}((c - v)^2)\right)$$



$$f = \text{NoiseFlood}\left(\text{ModSwitch}((c - v)^2)\right)$$

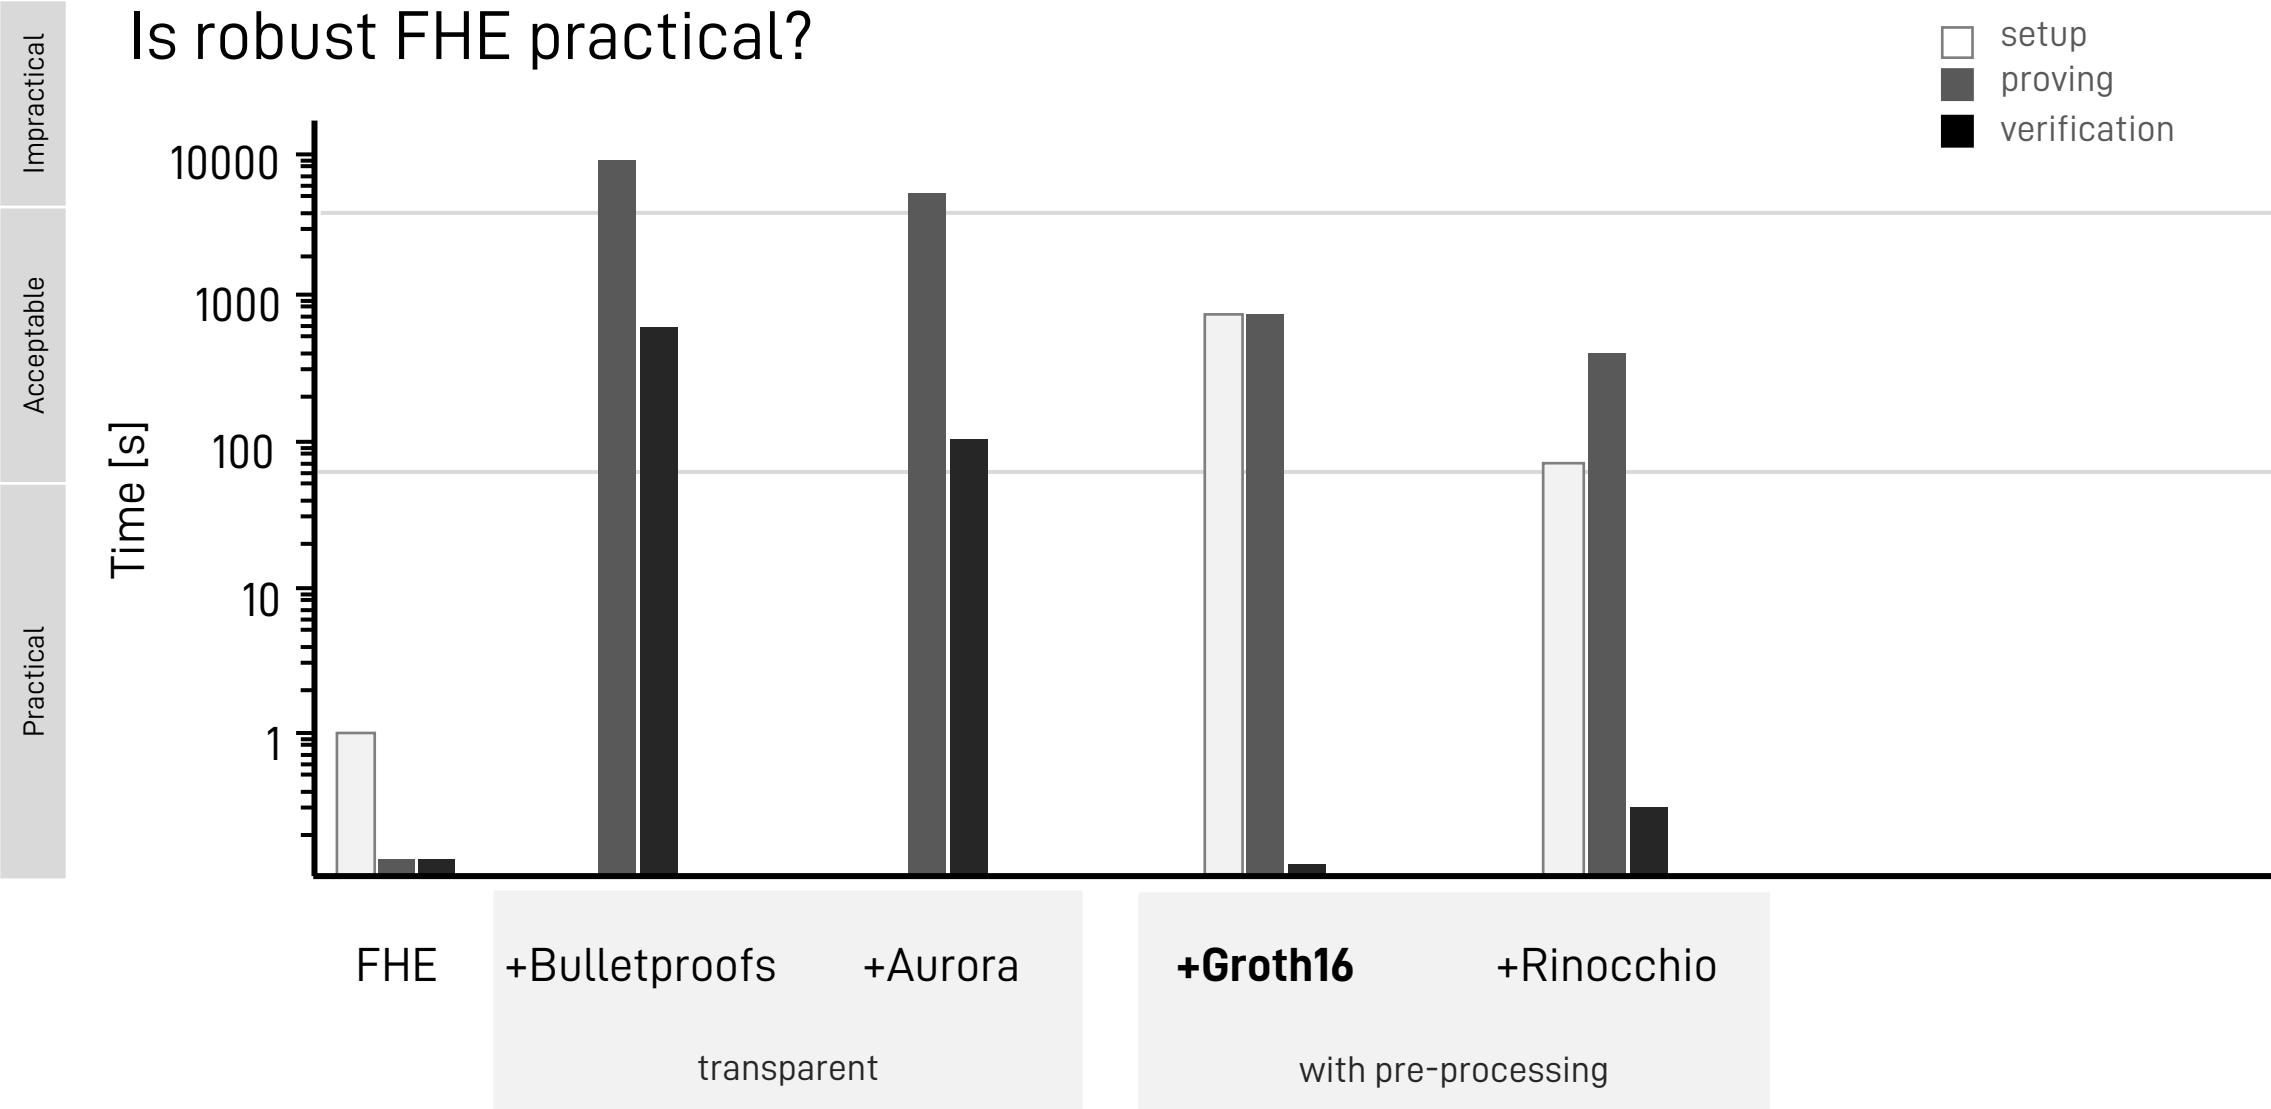


# Is robust FHE practical?



BGV, N = 8192,  $\log q = 183$ ,  $\log t = 30$

# Is robust FHE practical?

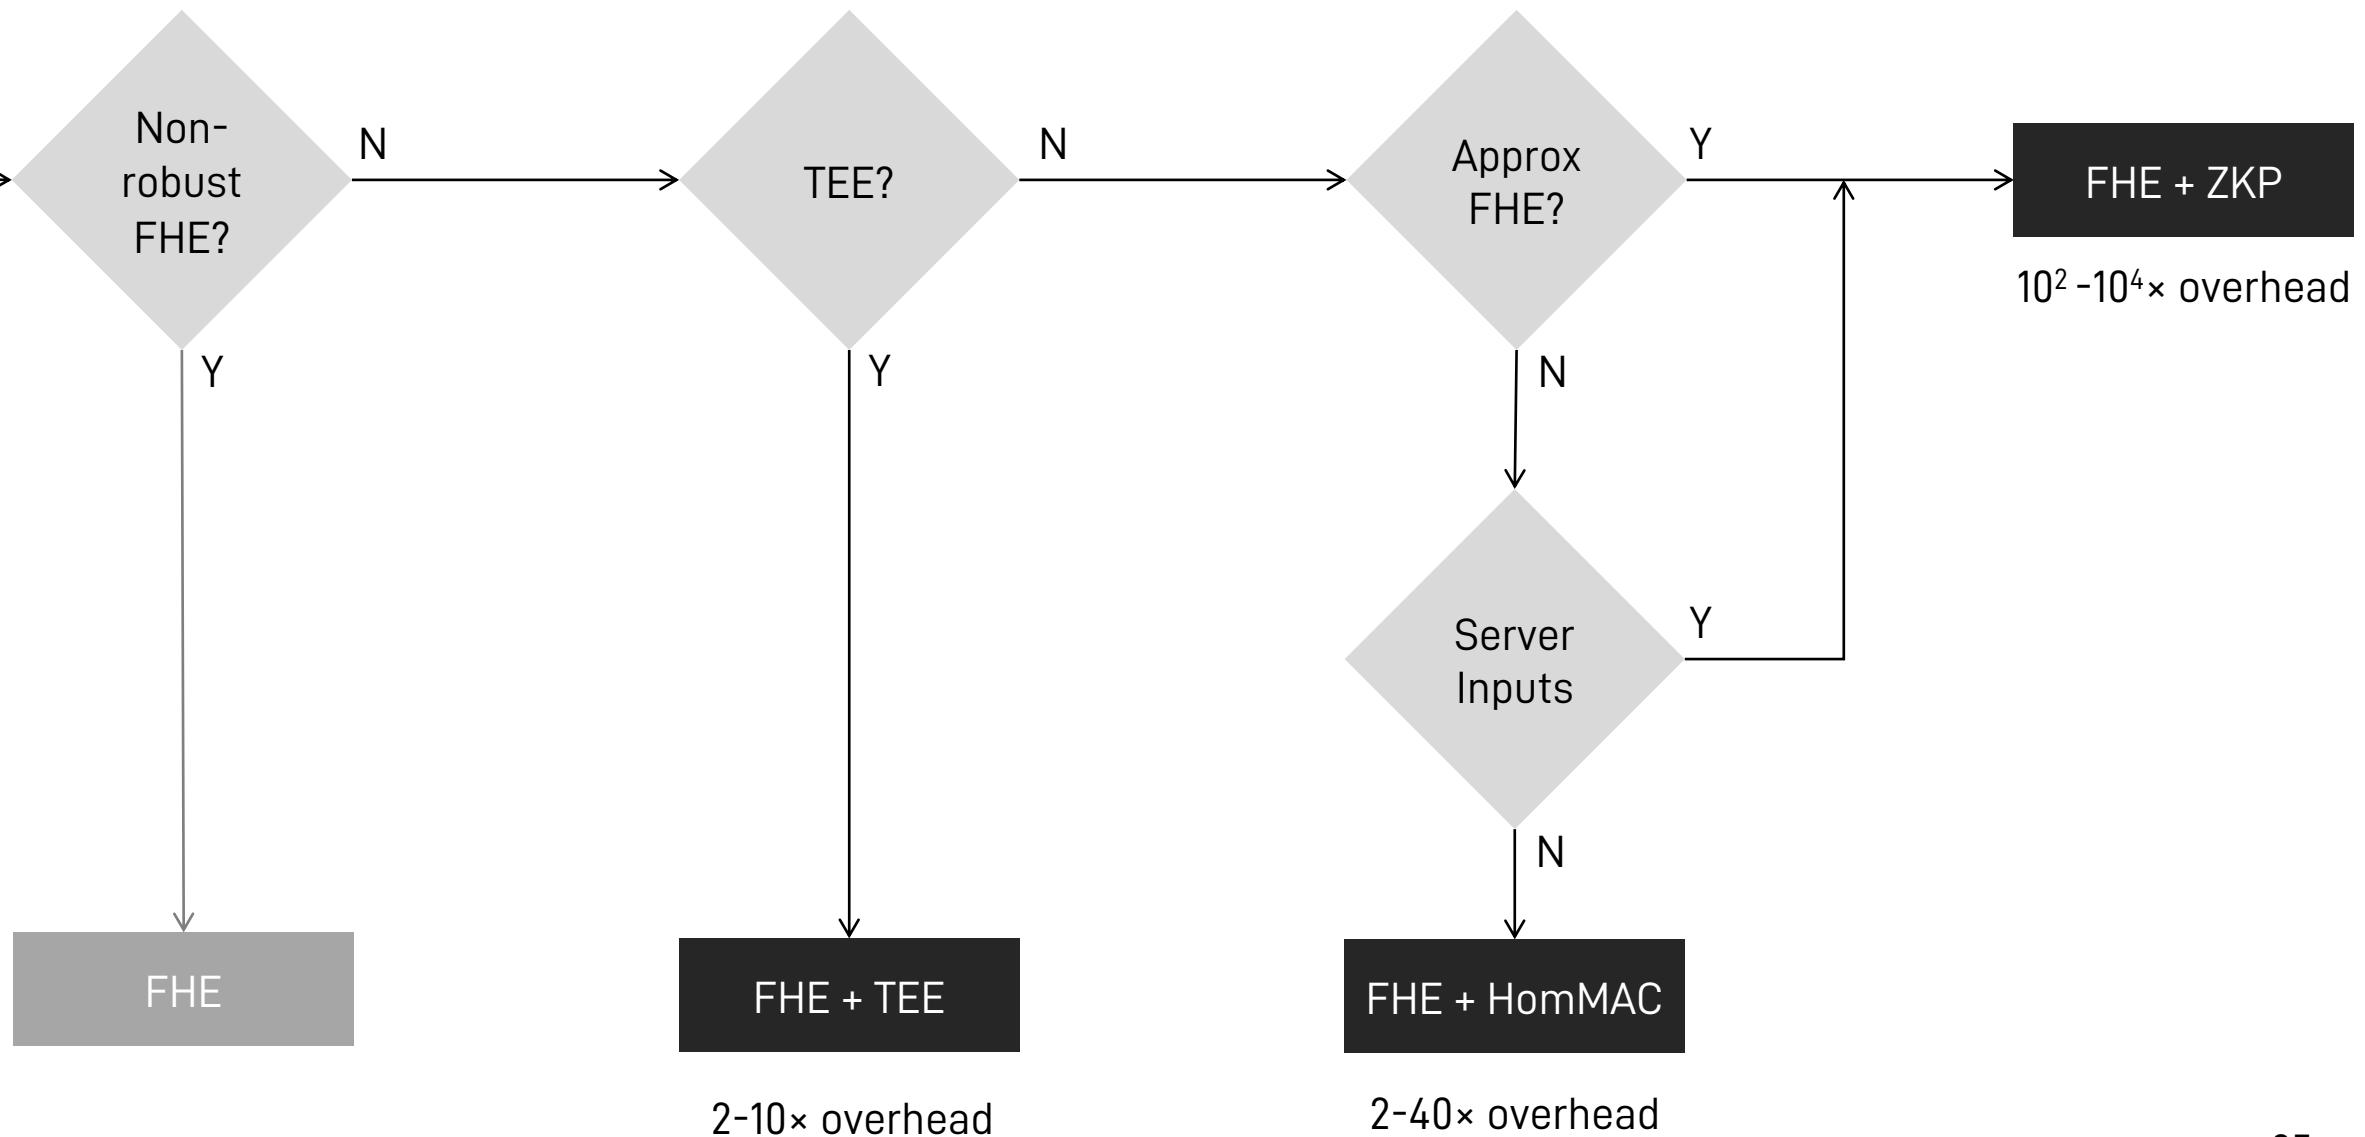


BGV, N = 8192,  $\log q = 183$ ,  $\log t = 30$

# Is robust FHE practical?

| TEE            | Field-ZKP    |                |             | Ring-ZKP   |
|----------------|--------------|----------------|-------------|------------|
|                | Transparent  | Pre-processing |             |            |
| SGX            | Bulletproofs | Aurora         | Groth16     | Rinocchio  |
| Small circuits | Practical    | Impractical    | Impractical | Acceptable |
| Big circuits   | Acceptable   | Impractical    | Impractical | Acceptable |

# Take-aways for practitioners



**This work:**

- Unified definitions for robust FHE
- Theoretic & concrete instantiation
- Improvements to FHE-in-TEE & ringSNARK



MarbleHE/FHE-in-TEE



MarbleHE/ringSNARK

[arXiv.org/abs/2301.07041](https://arxiv.org/abs/2301.07041)**Future work:**

- Stronger & application-specific notions
- More FHE-friendly proof systems
- Tooling & integration with compilers

**Verifiable Fully Homomorphic Encryption**Alexander Viand\*, Christian Knabenhans\*, Anwar Hithnawi  
ETH Zurich

*Abstract*—Fully Homomorphic Encryption (FHE) is seeing increasing real-world deployment to protect data in use by allowing computation over encrypted data. However, the same malleability that enables homomorphic computations also raises security issues, which have so far been mostly overlooked. Security and privacy has obvious implications for confidentiality:

**Honest-but-Curious Assumption.** Historically, the FHE research community has extensively made use of the assumption that the server running an FHE application would be honest-but-curious, rather than actively malicious [13]–[16]. This assumption may be reasonable in some deployment scenarios (e.g., when FHE is used only to ensure regulatory compliance or when dealing with trusted institutions cooperating on their own data). However, the necessity to make such an assumption to this extent is very limiting to the scope